

LAR and LASSO

Alexander McLain

LASSO and RR with the Bodyfat Data

This example will use the same bodyfat data from the MMST textbook.

```
bf_dat <- read.csv("bodyfat2.csv")
bf_df <- data.frame(bf_dat)
```

Let's recall the fit of the full model

```
bf_mod <- lm(bodyfat~age + weight + height + neck + chest +
             abdomen + hip + thigh + knee + ankle + biceps +
             forearm + wrist,data = bf_df)
round(coef(summary(bf_mod)),4)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-21.3532	22.1862	-0.9625	0.3368
age	0.0646	0.0322	2.0058	0.0460
weight	-0.0964	0.0618	-1.5584	0.1205
height	-0.0439	0.1787	-0.2459	0.8060
neck	-0.4755	0.2356	-2.0184	0.0447
chest	-0.0172	0.1032	-0.1665	0.8679
abdomen	0.9550	0.0902	10.5917	0.0000
hip	-0.1886	0.1448	-1.3025	0.1940
thigh	0.2483	0.1462	1.6991	0.0906
knee	0.0139	0.2477	0.0563	0.9552
ankle	0.1779	0.2226	0.7991	0.4251
biceps	0.1823	0.1725	1.0568	0.2917
forearm	0.4557	0.1993	2.2867	0.0231
wrist	-1.6545	0.5332	-3.1032	0.0021

Now we'll take a look a

```
library(MASS)
```

```
##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##      select
```

```
library(leaps)
leaps <- regsubsets(bodyfat~ age + weight + height + neck +
                   chest + abdomen + hip + thigh + knee +
                   ankle + biceps + forearm + wrist, data =
```

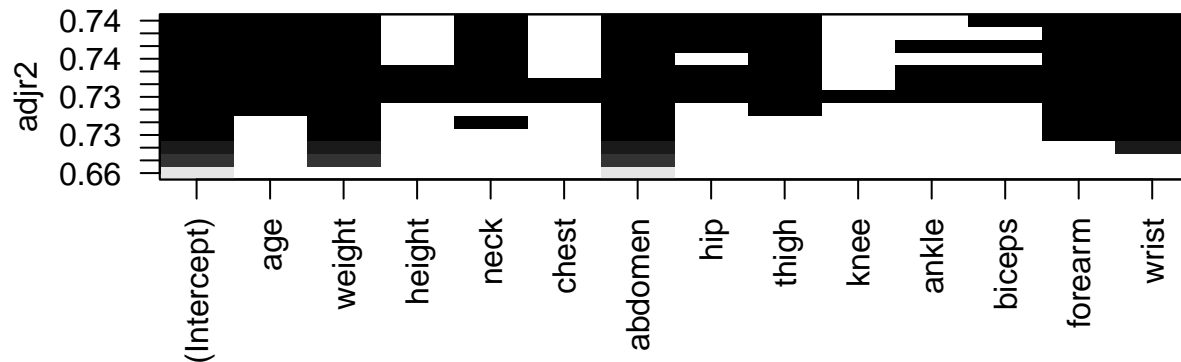
```

bf_df, method="exhaustive", nvmax=13)
summary(leaps)

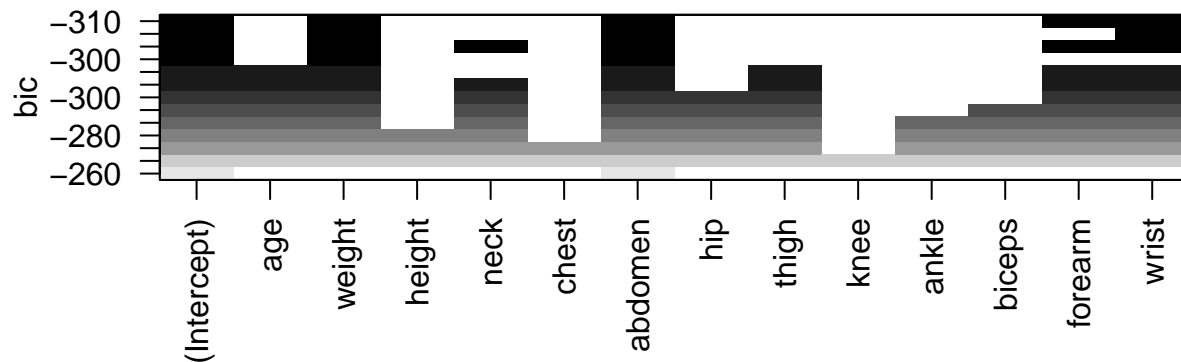
## Subset selection object
## Call: regsubsets.formula(bodyfat ~ age + weight + height + neck + chest +
##   abdomen + hip + thigh + knee + ankle + biceps + forearm +
##   wrist, data = bf_df, method = "exhaustive", nvmax = 13)
## 13 Variables (and intercept)
##      Forced in Forced out
## age          FALSE      FALSE
## weight        FALSE      FALSE
## height        FALSE      FALSE
## neck          FALSE      FALSE
## chest         FALSE      FALSE
## abdomen       FALSE      FALSE
## hip          FALSE      FALSE
## thigh        FALSE      FALSE
## knee         FALSE      FALSE
## ankle        FALSE      FALSE
## biceps       FALSE      FALSE
## forearm      FALSE      FALSE
## wrist        FALSE      FALSE
## 1 subsets of each size up to 13
## Selection Algorithm: exhaustive
##      age weight height neck chest abdomen hip thigh knee ankle biceps
## 1 ( 1 ) " " " " " " " " "*" " " " " " " " "
## 2 ( 1 ) " " "*" " " " " " " "*" " " " " " " " "
## 3 ( 1 ) " " "*" " " " " " " "*" " " " " " " " "
## 4 ( 1 ) " " "*" " " " " " " "*" " " " " " " " "
## 5 ( 1 ) " " "*" " " " "*" " " "*" " " " " " " " "
## 6 ( 1 ) "*" "*" " " " " " " "*" " " "*" " " " " " "
## 7 ( 1 ) "*" "*" " " " "*" " " "*" " " "*" " " " " " "
## 8 ( 1 ) "*" "*" " " " "*" " " "*" "*" "*" " " " " " "
## 9 ( 1 ) "*" "*" " " " "*" " " "*" "*" "*" " " " " "*"
## 10 ( 1 ) "*" "*" " " " "*" " " "*" "*" "*" " " "*" "*"
## 11 ( 1 ) "*" "*" "*" " "*" " " "*" "*" "*" " " "*" "*"
## 12 ( 1 ) "*" "*" "*" " "*" "*" "*" "*" "*" " " "*" "*"
## 13 ( 1 ) "*" "*" "*" " "*" "*" "*" "*" "*" "*" "*" "*"
##      forearm wrist
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " "*"
## 4 ( 1 ) "*" "*"
## 5 ( 1 ) "*" "*"
## 6 ( 1 ) "*" "*"
## 7 ( 1 ) "*" "*"
## 8 ( 1 ) "*" "*"
## 9 ( 1 ) "*" "*"
## 10 ( 1 ) "*" "*"
## 11 ( 1 ) "*" "*"
## 12 ( 1 ) "*" "*"
## 13 ( 1 ) "*" "*"

```

```
# plot a table of models showing variables in each model.
# models are ordered by the selection statistic.
plot(leaps,scale="adjr2") #plot options are scale=c("bic", "Cp", "adjr2", "r2").
```



```
plot(leaps,scale="bic")
```



Next we'll fit the data using the Least Angle Regression (LAR) algorithm.

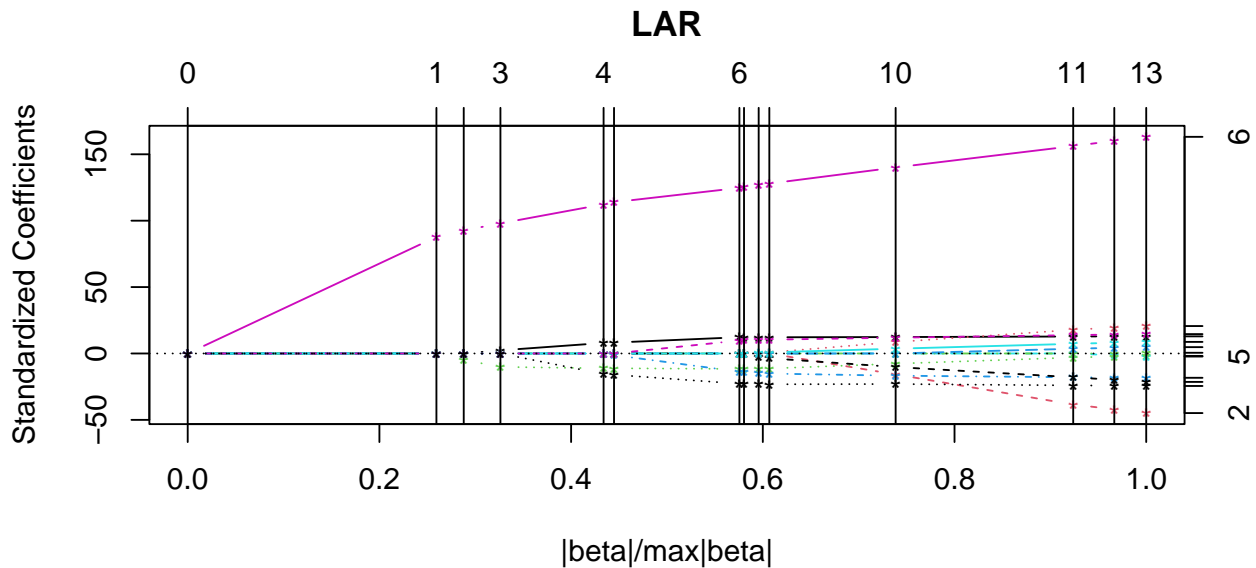
```
library(lars)
```

```
## Loaded lars 1.3
```

```
lars_mod <- lars( as.matrix(bf_df[,3:15]), bf_df$bodyfat,
  type = "lar")
print(lars_mod)
```

```
##
## Call:
## lars(x = as.matrix(bf_df[, 3:15]), y = bf_df$bodyfat, type = "lar")
## R-squared: 0.749
## Sequence of LAR moves:
##      abdomen height age wrist neck forearm hip biceps thigh weight ankle chest
## Var      6      3  1    13    4      12  7    11    8      2    10    5
## Step      1      2  3     4     5       6  7     8    9     10   11   12
##      knee
## Var      9
## Step    13
```

```
plot(lars_mod)
```



The last fit of the LAR procedure is equivalent to the OLS estimates:

```
cbind(round(lars_mod$beta[14,],4),round(bf_mod$coefficients[-1],4))
```

age	0.0646	0.0646
weight	-0.0964	-0.0964
height	-0.0439	-0.0439
neck	-0.4755	-0.4755
chest	-0.0172	-0.0172
abdomen	0.9550	0.9550
hip	-0.1886	-0.1886
thigh	0.2483	0.2483
knee	0.0139	0.0139
ankle	0.1779	0.1779
biceps	0.1823	0.1823
forearm	0.4557	0.4557
wrist	-1.6545	-1.6545

Let's see the lasso and ridge procedures.

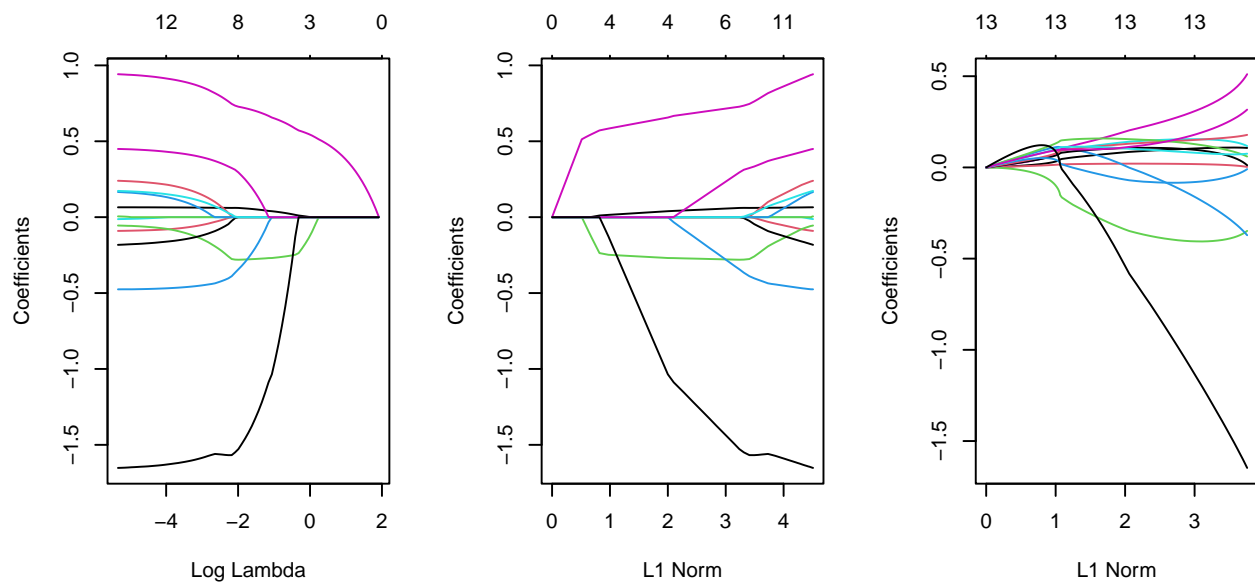
```
library(glmnet)
##Here we fit the LASSO model to our data. We'll look at this for the
##default nlambda = 100, and lambda.min.ratio = 0.0001.
names(bf_df)

## [1] "density" "bodyfat" "age"      "weight"  "height"  "neck"    "chest"
## [8] "abdomen" "hip"      "thigh"   "knee"    "ankle"   "biceps"  "forearm"
## [15] "wrist"

lasso_mod <- glmnet(as.matrix(bf_df[,3:15]), bf_df$bodyfat,
                    family="gaussian",alpha=1)
summary(lasso_mod)
```

	Length	Class	Mode
a0	79	-none-	numeric
beta	1027	dgCMatrix	S4
df	79	-none-	numeric
dim	2	-none-	numeric
lambda	79	-none-	numeric
dev.ratio	79	-none-	numeric
nulldev	1	-none-	numeric
npasses	1	-none-	numeric
jerr	1	-none-	numeric
offset	1	-none-	logical
call	5	-none-	call
nobs	1	-none-	numeric

```
ridge_mod <- glmnet(as.matrix(bf_df[,3:15]), bf_df$bodyfat,
  family="gaussian", alpha=0)
par(mfrow=c(1,3))
plot(lasso_mod,xvar="lambda",label=TRUE)
plot(lasso_mod,xvar="norm",label=TRUE)
###Not lets fit the ridge regression model.
plot(ridge_mod,xvar="norm",label=TRUE)
```



```
round(lasso_mod$beta[,32],3) #log(lambda)=-1
```

```
##      age  weight  height    neck  chest abdomen    hip  thigh  knee  ankle
##    0.037  0.000 -0.267   0.000   0.000   0.649   0.000   0.000   0.000   0.000
## biceps forearm  wrist
##    0.000  0.000 -0.939
```

```
round(lasso_mod$beta[,43],3) #log(lambda)=-2
```

```
##      age  weight  height    neck  chest abdomen    hip  thigh  knee  ankle
##    0.061  0.000 -0.280  -0.343   0.000   0.728   0.000   0.000   0.000   0.000
## biceps forearm  wrist
##    0.000  0.295 -1.528
```

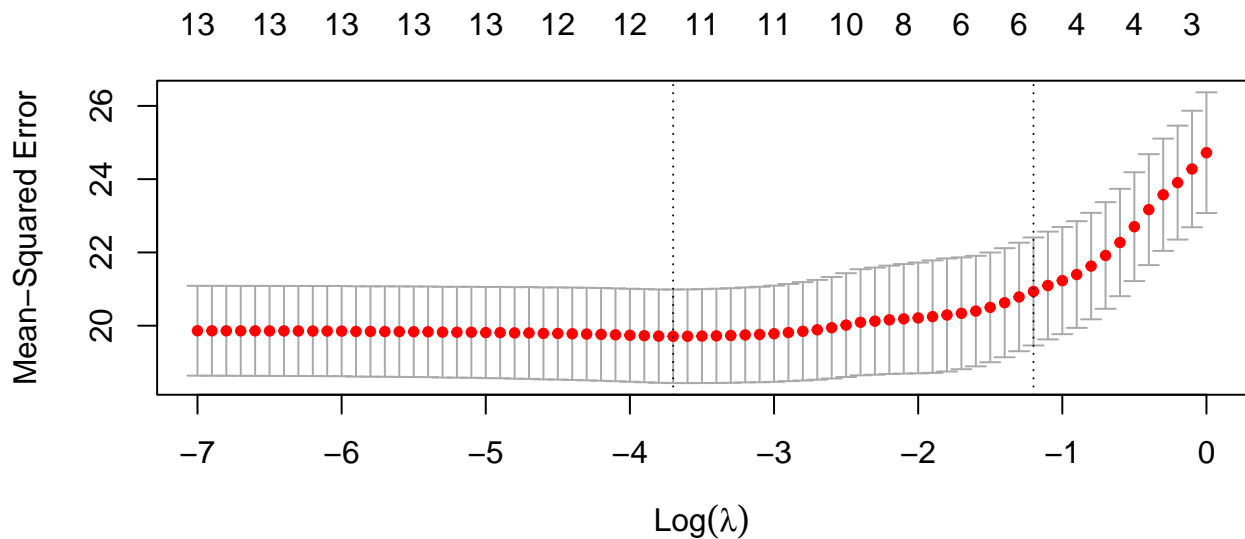
```
round(lasso_mod$beta[,65],4) #log(lambda)=-4
```

```
##    age  weight  height    neck  chest abdomen    hip  thigh    knee  ankle
## 0.0644 -0.0814 -0.0752 -0.4692 -0.0015  0.9152 -0.1624  0.2158  0.0000  0.1343
## biceps forearm  wrist
## 0.1528  0.4321 -1.6312
```

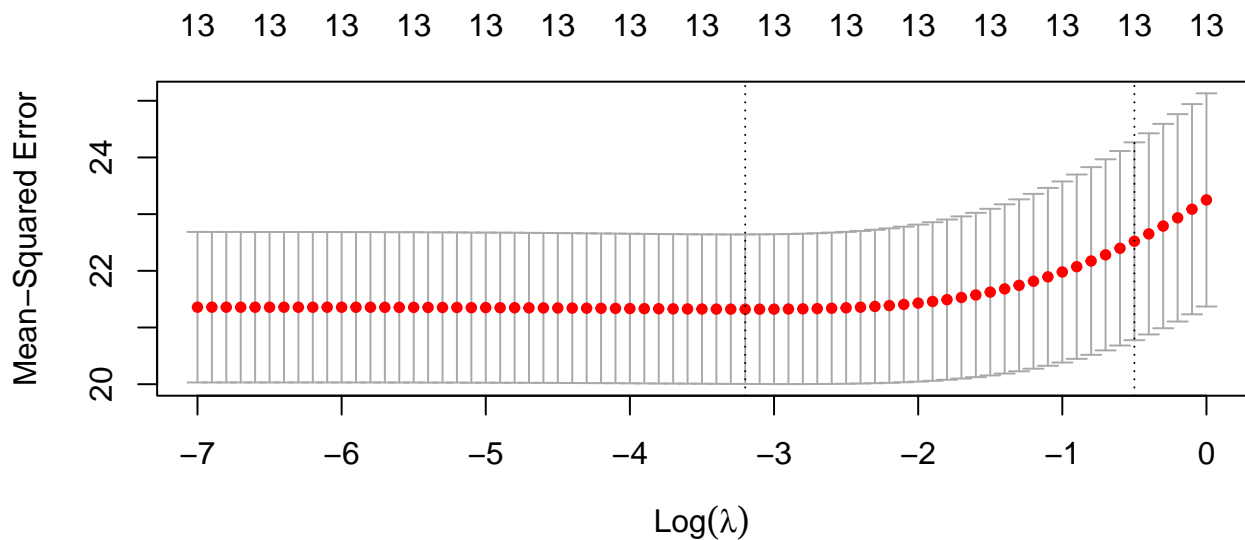
Which value of λ should we use? To determine that we can use cross-validation.

```
lasso_cv_mod <- cv.glmnet(as.matrix(bf_df[,3:15]), bf_df$bodyfat,
                          family="gaussian", alpha=1, lambda = exp( seq(-7, 0, 0.1)) )
ridge_cv_mod <- cv.glmnet(as.matrix(bf_df[,3:15]), bf_df$bodyfat,
                          family="gaussian", alpha=0, lambda = exp( seq(-7, 0, 0.1)) )
```

```
plot(lasso_cv_mod)
```



```
plot(ridge_cv_mod)
```



```
lasso_coef_1se <- as.matrix( coef.glmnet( lasso_cv_mod, s = c("lambda.1se")) )
lasso_coef_min <- as.matrix( coef.glmnet( lasso_cv_mod, s = c("lambda.min")) )
ridge_coef_1se <- as.matrix( coef.glmnet( ridge_cv_mod, s = c("lambda.1se")) )
```

```
ridge_coef_min <- as.matrix( coef.glmnet( ridge_cv_mod, s = c("lambda.min")) )
coef <- data.frame( lasso_coef_1se, lasso_coef_min, ridge_coef_1se, ridge_coef_min)
colnames(coef) <- c("lasso_1se", "lasso_min", "ridge_1se", "ridge_min")
round(coef,3)
```

	lasso_1se	lasso_min	ridge_1se	ridge_min
(Intercept)	-3.978	-16.301	-2.042	-14.687
age	0.043	0.064	0.106	0.073
weight	0.000	-0.075	0.000	-0.071
height	-0.271	-0.090	-0.338	-0.105
neck	-0.057	-0.465	-0.386	-0.485
chest	0.000	0.000	0.111	-0.005
abdomen	0.672	0.902	0.532	0.891
hip	0.000	-0.153	0.001	-0.175
thigh	0.000	0.201	0.181	0.237
knee	0.000	0.000	0.054	0.011
ankle	0.000	0.115	-0.001	0.147
biceps	0.000	0.142	0.077	0.158
forearm	0.025	0.423	0.328	0.444
wrist	-1.123	-1.620	-1.671	-1.700

Grouped LASSO

This will be a quick example that uses grouped lasso with the `gglasso` package.

```
library(gglasso)
`?`(gglasso)
```

Fits the regularization paths for group-lasso penalized learning problems

Description:

Fits regularization paths for group-lasso penalized learning problems at a sequence of regularization parameters λ .

Usage:

```
gglasso(
  x,
  y,
  group = NULL,
  loss = c("ls", "logit", "sqsvm", "hsvm", "wls"),
  nlambdas = 100,
  lambda.factor = ifelse(nobs < nvars, 0.05, 0.001),
  lambda = NULL,
  pf = sqrt(bs),
  weight = NULL,
  dfmax = as.integer(max(group)) + 1,
  pmax = min(dfmax * 1.2, as.integer(max(group))),
  eps = 1e-08,
  maxit = 3e+08,
```

```

    delta,
    intercept = TRUE
)

```

Arguments:

x: matrix of predictors, of dimension $n \times p$; each row is an observation vector.

y: response variable. This argument should be quantitative for regression (least squares), and a two-level factor for classification (logistic model, huberized SVM, squared SVM).

group: a vector of consecutive integers describing the grouping of the coefficients (see example below).

loss: a character string specifying the loss function to use, valid options are:

- "ls" least squares loss (regression),
- "logit" logistic loss (classification).
- "hsvm" Huberized squared hinge loss (classification),
- "sqsvm" Squared hinge loss (classification),

Default is "ls".

nlambda: the number of 'lambda' values - default is 100.

lambda.factor: the factor for getting the minimal lambda in 'lambda' sequence, where $\text{'min(lambda)'} = \text{'lambda.factor'} \times \text{'max(lambda)'}$. 'max(lambda)' is the smallest value of 'lambda' for which all coefficients are zero. The default depends on the relationship between n (the number of rows in the matrix of predictors) and p (the number of predictors). If $n \geq p$, the default is '0.001', close to zero. If $n < p$, the default is '0.05'. A very small value of 'lambda.factor' will lead to a saturated fit. It takes no effect if there is user-defined 'lambda' sequence.

lambda: a user supplied 'lambda' sequence. Typically, by leaving this option unspecified users can have the program compute its own 'lambda' sequence based on 'nlambda' and 'lambda.factor'. Supplying a value of 'lambda' overrides this. It is better to supply a decreasing sequence of 'lambda' values than a single (small) value, if not, the program will sort user-defined 'lambda' sequence in decreasing order automatically.

pf: penalty factor, a vector in length of b_n (b_n is the total number of groups). Separate penalty weights can be applied to each group of beta's to allow differential shrinkage. Can be 0 for some groups, which implies no shrinkage, and results in

that group always being included in the model. Default value for each entry is the square-root of the corresponding size of each group.

weight: a nxn observation weight matrix in the where n is the number of observations. Only used if 'loss='wls'' is specified. Note that cross-validation is NOT IMPLEMENTED for 'loss='wls''.

dfmax: limit the maximum number of groups in the model. Useful for very large 'bs' (group size), if a partial path is desired. Default is 'bs+1'.

pmax: limit the maximum number of groups ever to be nonzero. For example once a group enters the model, no matter how many times it exits or re-enters model through the path, it will be counted only once. Default is 'min(dfmax*1.2,bs)'.

eps: convergence termination tolerance. Defaults value is '1e-8'.

maxit: maximum number of outer-loop iterations allowed at fixed lambda value. Default is 3e8. If models do not converge, consider increasing 'maxit'.

delta: the parameter delta in '"hsvm"' (Huberized squared hinge loss). Default is 1.

intercept: Whether to include intercept in the model. Default is TRUE.

```
data(bardet)
`?`(bardet)
```

Simplified gene expression data from Scheetz et al. (2006)

Description:

Gene expression data (20 genes for 120 samples) from the microarray experiments of mammalian eye tissue samples of Scheetz et al. (2006).

Details:

This data set contains 120 samples with 100 predictors (expanded from 20 genes using 5 basis B-splines, as described in Yang, Y. and Zou, H. (2015)).

Value:

A list with the following elements:

x: a [120 x 100] matrix (expanded from a [120 x 20] matrix) giving the expression levels of 20 filtered genes for the 120 samples. Each row corresponds to a subject, each 5 consecutive columns to a grouped gene.

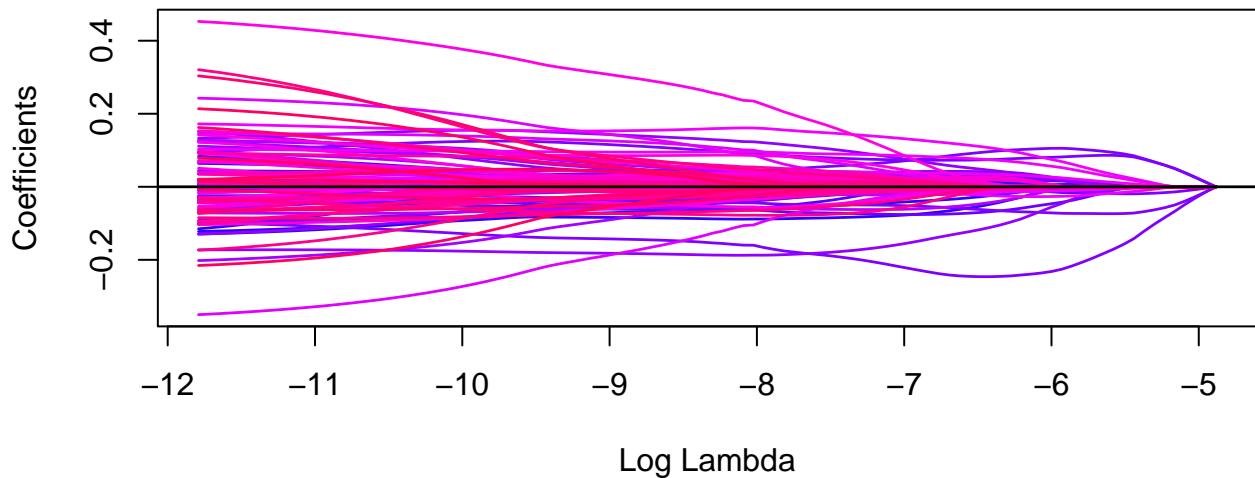
y: a numeric vector of length 120 giving expression level of

gene TRIM32, which causes Bardet-Biedl syndrome.

Now we'll fit the group lasso to the bardet data

```
# define group index
group1 <- rep(1:20, each = 5)

# fit group lasso penalized least squares
m1 <- gglasso(x = bardet$x, y = bardet$y, group = group1, loss = "ls")
plot(m1)
```



Now we'll do some CV to properly choose the value of λ .

```
cv <- cv.gglasso(x = bardet$x, y = bardet$y, group = group1, loss = "ls", pred.loss = "L2",
  lambda.factor = 0.05, nfolds = 5)
cv$lambda.min

## [1] 0.0004681541

cv$lambda.1se

## [1] 0.007575771

cbind(c(0, group1), c(coef(cv, s = "lambda.1se")), c(coef(cv, s = "lambda.min")))
```

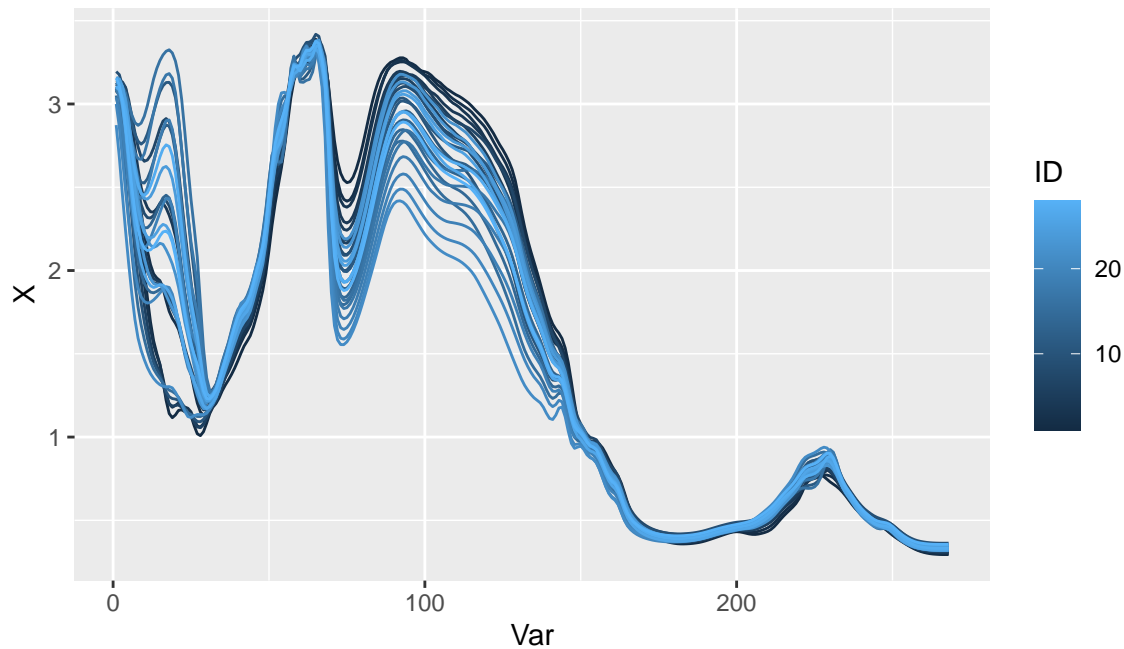
0	8.390844	8.2173271
1	0.000000	-0.0137397
1	0.000000	-0.0523389
1	0.000000	0.0356671
1	0.000000	0.0090829
1	0.000000	-0.0838165
2	0.000000	0.0002781
2	0.000000	0.0017139
2	0.000000	0.0012646
2	0.000000	0.0001063
2	0.000000	-0.0028337
3	0.000000	0.0220613
3	0.000000	-0.0254643
3	0.000000	-0.0007793
3	0.000000	0.0028207
3	0.000000	-0.0452385
4	0.000000	0.0046160

4	0.000000	0.0442182
4	0.000000	-0.0227353
4	0.000000	-0.0058258
4	0.000000	-0.0702611
5	0.000000	0.0508213
5	0.000000	0.0818889
5	0.000000	-0.0491862
5	0.000000	-0.0181138
5	0.000000	-0.1756574
6	0.000000	0.1123866
6	0.000000	0.0339147
6	0.000000	-0.0271549
6	0.000000	0.0083516
6	0.000000	-0.1898404
7	0.000000	-0.0013961
7	0.000000	0.0005563
7	0.000000	0.0006054
7	0.000000	0.0010068
7	0.000000	0.0014450
8	0.000000	0.0357903
8	0.000000	-0.0202496
8	0.000000	0.0076874
8	0.000000	0.0117014
8	0.000000	-0.0607913
9	0.000000	-0.0004775
9	0.000000	-0.0006332
9	0.000000	0.0004213
9	0.000000	0.0015303
9	0.000000	0.0006469
10	0.000000	-0.0639107
10	0.000000	0.0487546
10	0.000000	0.0760078
10	0.000000	0.0110616
10	0.000000	0.0166093
11	0.000000	-0.0126953
11	0.000000	0.0017876
11	0.000000	0.0898015
11	0.000000	0.1500386
11	0.000000	0.0413499
12	0.000000	0.0000000
12	0.000000	0.0000000
12	0.000000	0.0000000
12	0.000000	0.0000000
12	0.000000	0.0000000
13	0.000000	-0.0520235
13	0.000000	0.0262814
13	0.000000	-0.0005512
13	0.000000	0.0788385
13	0.000000	0.1842589
14	0.000000	-0.0261599
14	0.000000	0.0062349
14	0.000000	0.0259086
14	0.000000	0.0338999

14	0.000000	0.0037641
15	0.000000	-0.0095728
15	0.000000	-0.0039897
15	0.000000	0.0310495
15	0.000000	-0.0030660
15	0.000000	0.0132002
16	0.000000	-0.0182053
16	0.000000	0.0172314
16	0.000000	0.0460860
16	0.000000	-0.0025930
16	0.000000	0.0182979
17	0.000000	0.0051751
17	0.000000	-0.0138591
17	0.000000	0.0070085
17	0.000000	0.0067323
17	0.000000	-0.0220390
18	0.000000	0.0094923
18	0.000000	-0.0584266
18	0.000000	0.0092960
18	0.000000	0.0365321
18	0.000000	-0.0790914
19	0.000000	0.0048644
19	0.000000	-0.0012055
19	0.000000	-0.0018575
19	0.000000	0.0039680
19	0.000000	-0.0068155
20	0.000000	0.0000000
20	0.000000	0.0000000
20	0.000000	0.0000000
20	0.000000	0.0000000
20	0.000000	0.0000000

Re-analysis of the PET data

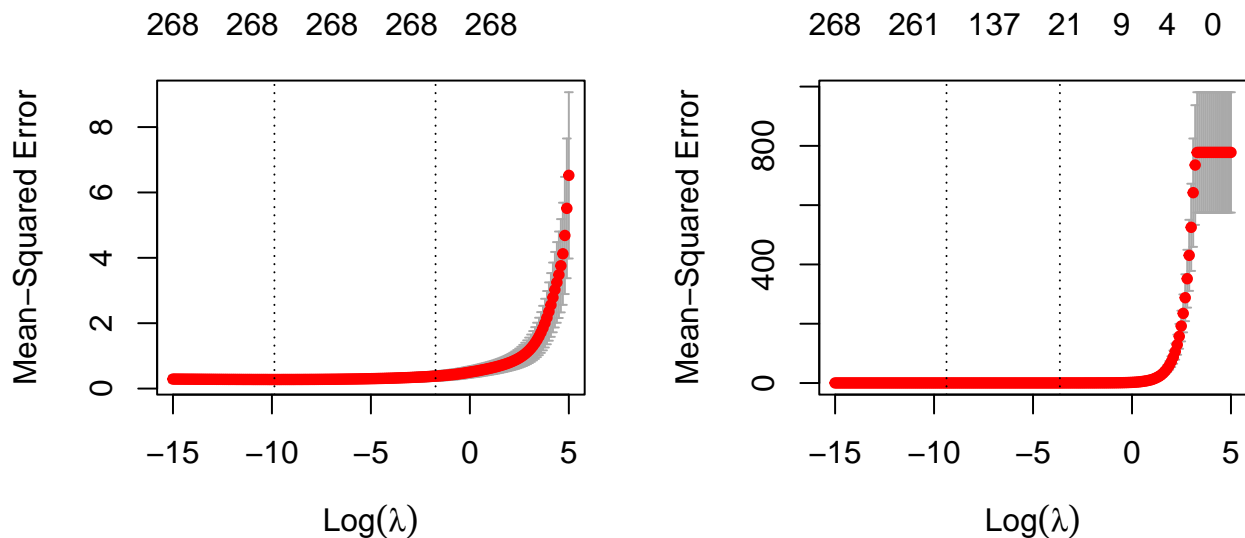
[1] 28 269



Now lets use `cv.glmnet` to fit ridge regression and lasso models.

```
ridge_cv_mod <- cv.glmnet(pet_mat[, 1:268], pet_mat[, 269], family = "gaussian",
  alpha = 0, lambda = exp(seq(-15, 5, length.out = 200)), grouped = FALSE)
```

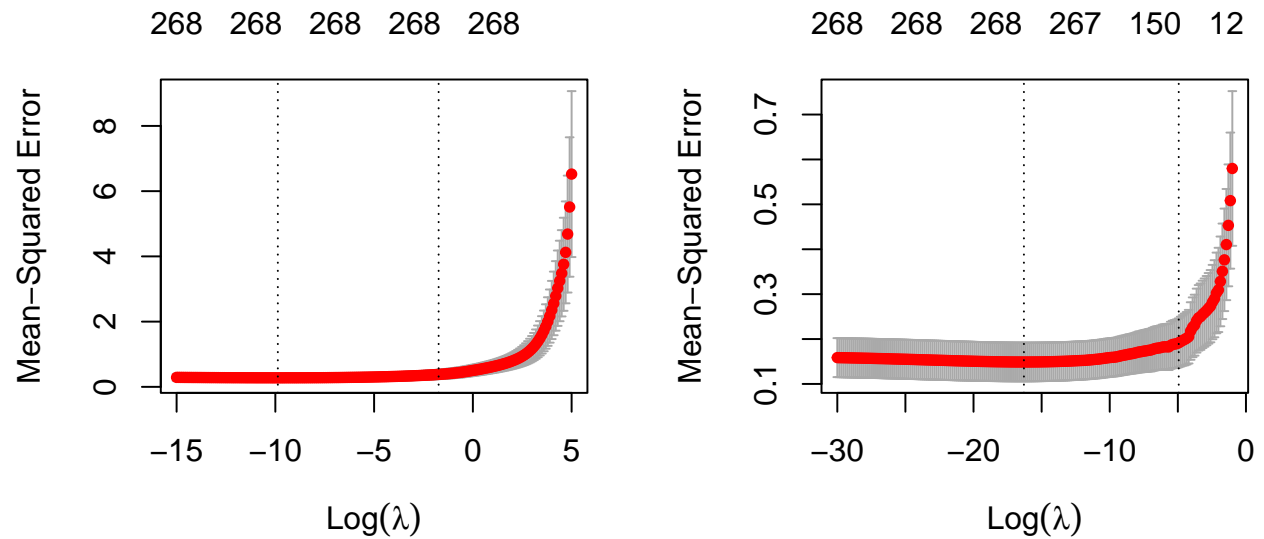
```
lasso_cv_mod <- cv.glmnet(pet_mat[, 1:268], pet_mat[, 269], family = "gaussian",
  alpha = 1, lambda = exp(seq(-15, 5, length.out = 200)), grouped = FALSE)
par(mfrow = c(1, 2))
plot(ridge_cv_mod)
plot(lasso_cv_mod)
```



We can see that the best lasso model had the smallest possible value of λ . As a result, we'll fit again with some smaller values.

```
lasso_cv_mod <- cv.glmnet(pet_mat[, 1:268], pet_mat[, 269], family = "gaussian",
  alpha = 1, lambda = exp(seq(-30, -1, length.out = 200)), grouped = FALSE)
par(mfrow = c(1, 2))
```

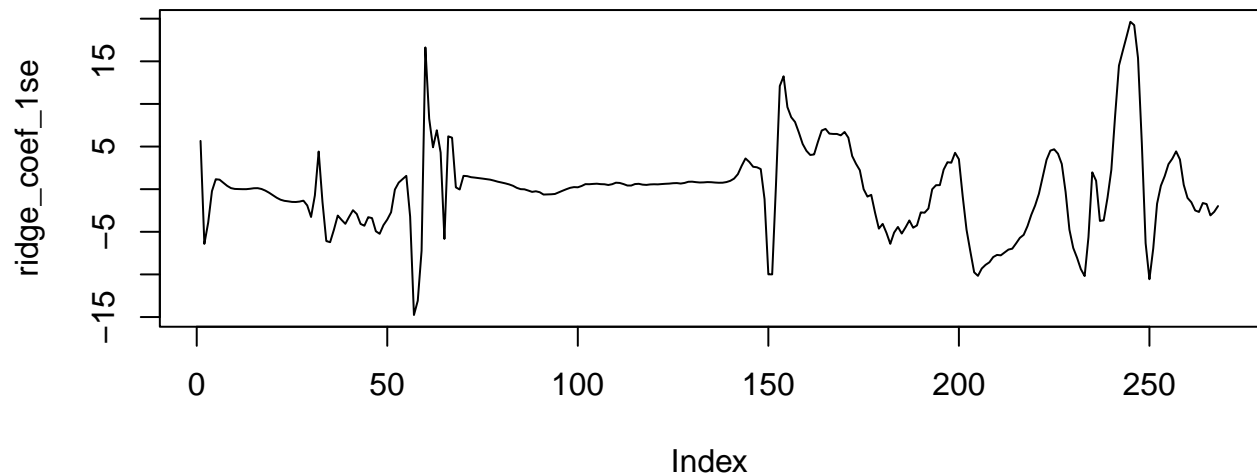
```
plot(ridge_cv_mod)
plot(lasso_cv_mod)
```



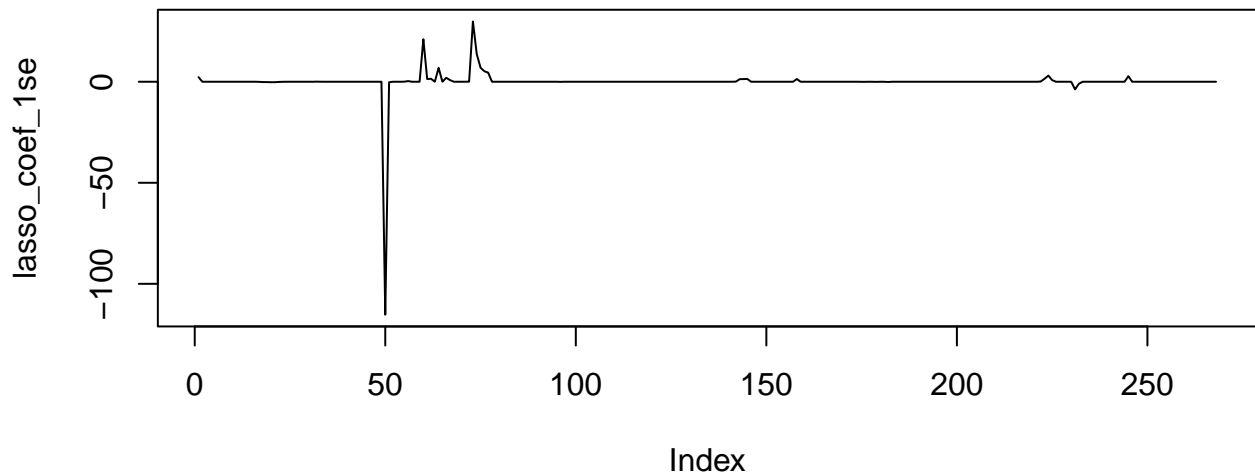
Let's look at the results:

```
ridge_coef_1se <- coef.glmnet( ridge_cv_mod,
                               s = c("lambda.1se"))[-1]
lasso_coef_1se <- coef.glmnet( lasso_cv_mod,
                               s = c("lambda.1se"))[-1]

plot(ridge_coef_1se, type = "l")
```



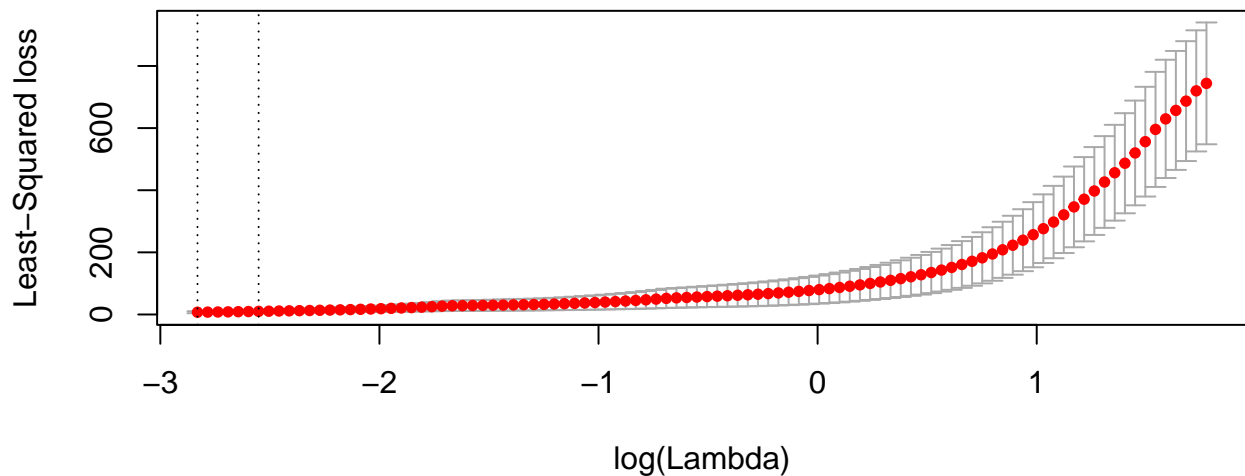
```
plot(lasso_coef_1se, type = "l")
```



While there doesn't appear to be any grouping or similar effect by variable number, it does make sense if we think about the data. As a result, let's create a grouping variable and analyze this with the grouped lasso.

```
pet_group <- rep(1:14, each = 20)[1:268]
gglasso_cv <- cv.gglasso(x = pet_mat[, 1:268], y = pet_mat[, 269], group = pet_group,
  loss = "ls", pred.loss = "L2", lambda.factor = 0.01, nfolds = 5)
par(mfrow = c(1, 1))
plot(gglasso_cv)
```

-2.83043616 -1.85358188 -0.87672760 0.05360981 0.93743035



Let's look at the results:

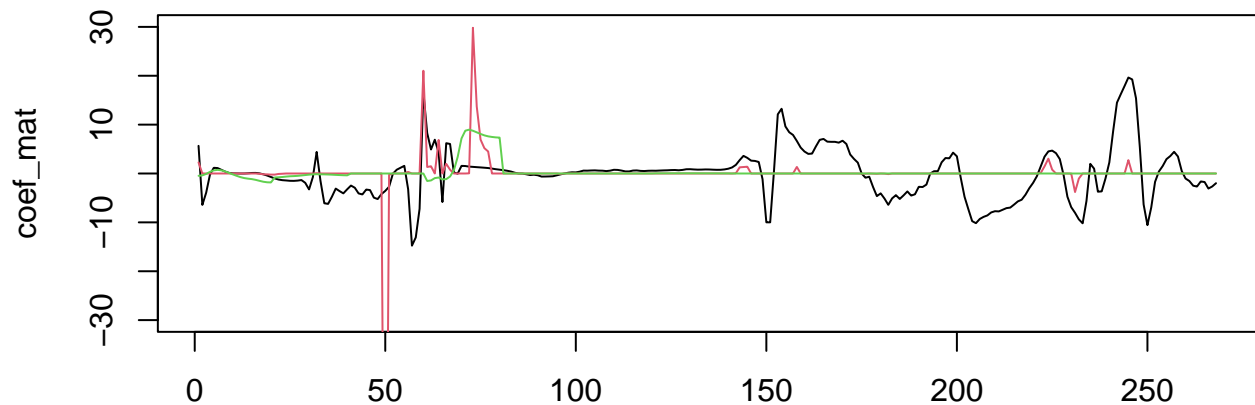
```
gglasso_coef_1se <- coef(gglasso_cv,
  s="lambda.1se")[-1]
coef_mat <- cbind(ridge_coef_1se,
  lasso_coef_1se,
  gglasso_coef_1se)
table(lasso_coef_1se==0)
```

FALSE	TRUE
55	213

```
table(gglasso_coef_1se==0)
```

FALSE	TRUE
80	188

```
matplot(coef_mat, type = "l",
        ylim = c(-30,30), lty = 1)
```



Now let's look at this with the fused LASSO.

```
library(genlasso)
`?`(fusedlasso)
```

Compute the fused lasso solution path for a general graph, or a 1d or 2d grid

Description:

These functions produce the solution path for a general fused lasso problem. The 'fusedlasso' function takes either a penalty matrix or a graph object from the 'igraph' package. The 'fusedlassoid' and 'fusedlasso2d' functions are convenience functions that construct the penalty matrix over a 1d or 2d grid.

Usage:

```
fusedlasso(y, X, D, graph, gamma = 0, approx = FALSE, maxsteps = 2000,
           minlam = 0, rtol = 1e-07, btol = 1e-07, eps = 1e-4,
           verbose = FALSE)
fusedlassoid(y, pos, X, gamma = 0, approx = FALSE, maxsteps = 2000,
             minlam = 0, rtol = 1e-07, btol = 1e-07, eps = 1e-4,
             verbose = FALSE)
fusedlasso2d(y, X, dim1, dim2, gamma = 0, approx = FALSE, maxsteps = 2000,
             minlam = 0, rtol = 1e-07, btol = 1e-07, eps = 1e-4,
             verbose = FALSE)
```

Arguments:

y: a numeric response vector. Alternatively, for 'fusedlasso2d' with no matrix 'X' passed, 'y' can be a matrix (its

dimensions corresponding to the underlying 2d grid). Note that when 'y' is given as a vector in 'fusedlasso2d', with no 'X' passed, it should be in column major order.

pos: only for 'fusedlassoid', these are the optional positions of the positions in the 1d grid. If missing, the 1d grid is assumed to have unit spacing.

X: an optional matrix of predictor variables, with observations along the rows, and variables along the columns. If the passed 'X' has more columns than rows, then a warning is given, and a small ridge penalty is added to the generalized lasso criterion before the path is computed. If 'X' has less columns than rows, then its rank is not checked for efficiency, and (unlike the 'genasso' function) a ridge penalty is not automatically added if it is rank deficient. Therefore, a tall, rank deficient 'X' may cause errors.

D: only for 'fusedlasso', this is the penalty matrix, i.e., the oriented incidence matrix over the underlying graph (the orientation of each edge being arbitrary). Only one of 'D' or 'graph' needs to be specified.

graph: only for 'fusedlasso', this is the underlying graph as an 'igraph' object from the 'igraph' package. Only one of 'D' or 'graph' needs to be specified.

dim1: only for 'fusedlasso2d', this is the number of rows in the underlying 2d grid. If missing and 'y' is given as a matrix, it is assumed to be the number of rows of 'y'.

dim2: only for 'fusedlasso2d', this is the number of columns in the underlying 2d grid. If missing and 'y' is given as a matrix, it is assumed to be the number of columns of 'y'.

gamma: a numeric variable greater than or equal to 0, indicating the ratio of the two tuning parameters, one for the fusion penalty, and the other for the pure ℓ_1 penalty. Default is 0. See "Details" for more information.

approx: a logical variable indicating if the approximate solution path should be used (with no dual coordinates leaving the boundary). Default is 'FALSE'. Note that for the 1d fused lasso, with identity predictor matrix, this approximate path is the same as the exact solution path.

maxsteps: an integer specifying the maximum number of steps for the algorithm to take before termination. Default is 2000.

minlam: a numeric variable indicating the value of lambda at which the path should terminate. Default is 0.

rtol: a numeric variable giving the tolerance for determining the rank of a matrix: if a diagonal value in the R factor of a QR

decomposition is less than R , in absolute value, then it is considered zero. Hence making `rtol` larger means being less stringent with determination of matrix rank. In general, do not change this unless you know what you are getting into! Default is $1e-7$.

`btol`: a numeric variable giving the tolerance for accepting "late" hitting and leaving times: future hitting times and leaving times should always be less than the current knot in the path, but sometimes for numerical reasons they are larger; any computed hitting or leaving time larger than the current knot + `btol` is thrown away. Hence making `btol` larger means being less stringent with the determination of hitting and leaving times. Again, in general, do not change this unless you know what you are getting into! Default is $1e-7$.

`eps`: a numeric variable indicating the multiplier for the ridge penalty, in the case that 'X' is wide (more columns than rows). If numeric problems occur, make 'eps' larger. Default is $1e-4$.

`verbose`: a logical variable indicating if progress should be reported after each knot in the path.

Details:

The fused lasso estimate minimizes the criterion

$$\frac{1}{2} \sum_{i=1}^n (y_i - x_i^T \beta_i)^2 + \lambda \sum_{(i,j) \in E} |\beta_i - \beta_j| + \gamma \sum_{i=1}^p |\beta_i|,$$

where x_i is the i th row of the predictor matrix and E is the edge set of the underlying graph. The solution $\hat{\beta}$ is computed as a function of the regularization parameter λ , for a fixed value of γ . The default is to set $\gamma=0$, which corresponds to pure fusion of the coefficient vector β . A choice $\gamma>0$ introduces both sparsity and fusion in the coefficient vector, with a higher value placing more priority on sparsity.

If the predictor matrix is the identity, and the primal solution path β is desired at several levels of the ratio parameter γ , it is much more efficient to compute the solution path once with $\gamma=0$, and then use soft-thresholding via the 'softthresh' function.

Finally, for the image denoising problem, i.e., the fused lasso over a 2d grid with identity predictor matrix, it is easy to specify a huge graph with a seemingly small amount of data. For instance, running the 2d fused lasso (with identity predictor matrix) on an image at standard 1080p HD resolution yields a graph with over 2 million edges. Moreover, in image denoising problems-somewhat unlike most other applications of the fused lasso (and generalized lasso)-a solution is often desired near the dense end of the path ($\lambda=0$) as opposed to the regularized end

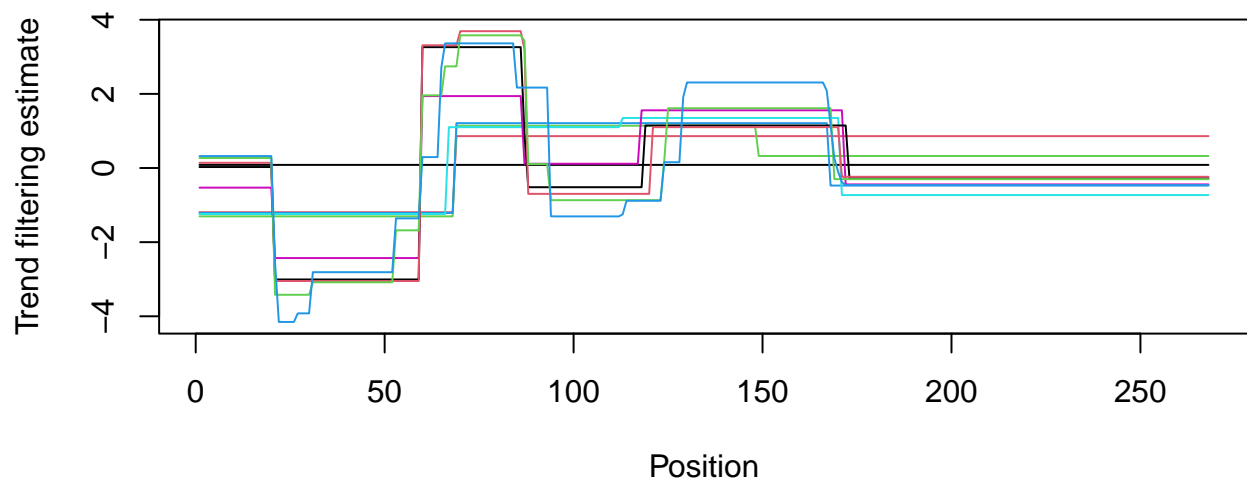
(lambda=Inf). The dual path algorithm implemented by the 'fusedlasso2d' function begins at the fully regularized end and works its way down to the dense end. For a problem with many edges (dual variables), if a solution at the dense is desired, then it must usually pass through a huge number knots in the path. Hence it is not advisable to run 'fusedlasso2d' on image denoising problems of large scale, as the dual solution path is computationally infeasible. It should be noted that a faster algorithm for the 2d fused lasso solution path (when the predictor matrix is the identity), which begins at the dense end of the path, is available in the 'flsa' package.

Since our X variables lie on a 1-dimensional line, we'll use the fusedlassoid version of the function.

```
fusedlasso <- fusedlassoid(y = pet_mat[, 269], pos = 1:268, X = pet_mat[, 1:268])

## Warning in trendfilter(y, pos, X, 0, approx, maxsteps, minlam, rtol, btol, :
## Adding a small ridge penalty (multiplier 0.0001), because X has more columns
## than rows.

par(mfrow = c(1, 1))
plot(fusedlasso)
```



This package doesn't have a built in CV function and there's really no way to manually do CV because this package doesn't have a way of manually selecting the λ values. Possibly other packages can do this.