

Multi-class Logistic, LDA and logistic lasso

ACM

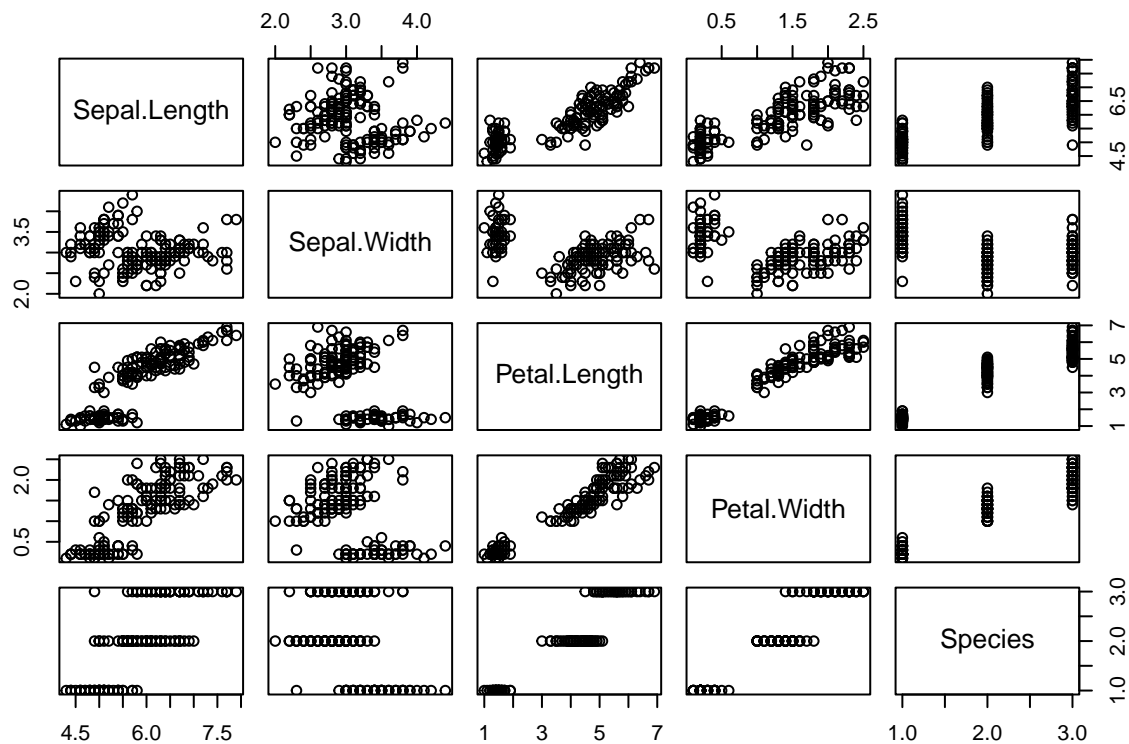
September 2023

Here we'll look at a quick example of multi-class LDA and logistic regression.


```
library(printr)
```

```
## Registered S3 method overwritten by 'printr':  
##   method      from  
##   knit_print.data.frame rmarkdown
```

```
library(ggplot2)  
# Load the Iris data  
data(iris)  
plot(iris)
```



```
plot(Petal.Width ~ Petal.Length, col=c("black", "red", "blue")[Species],  
     pch=(15:17)[Species], xlab="Petal Length (cm)",  
     ylab="Petal Width (cm)", data=iris)
```



```
head(iris)
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.0	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.0	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa

```
library(MASS)
set.seed(9823479)
train <- sample(1:150, 100)
table(iris$Species[train])
```

setosa	versicolor	virginica
31	35	34

```
z <- lda(Species ~ ., data = iris, subset = train)
z
```

Call:

```
lda(Species ~ ., data = iris, subset = train)
```

Prior probabilities of groups:

setosa	versicolor	virginica
0.31	0.35	0.34

Group means:

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
--------------	-------------	--------------	-------------

setosa	5.019355	3.412903	1.461290	0.2419355
versicolor	5.960000	2.825714	4.302857	1.3542857
virginica	6.567647	3.023529	5.555882	2.0411765

Coefficients of linear discriminants:

	LD1	LD2
Sepal.Length	0.3263448	0.2583307
Sepal.Width	2.3175531	-2.6202166
Petal.Length	-2.0326574	0.7725431
Petal.Width	-2.8992573	-2.7197269

Proportion of trace:

LD1	LD2
0.9911	0.0089

Now let's predict for the test data:

```
pred <- predict(z, iris[-train, ])$class
tbl <- table(pred, iris$Species[-train])
prop.table(tbl)
```

pred/	setosa	versicolor	virginica
setosa	0.38	0.0	0.00
versicolor	0.00	0.3	0.02
virginica	0.00	0.0	0.30

```
sum(diag(prop.table(tbl)))
```

```
## [1] 0.98
```

```
library(nnet)
`?`(multinom)
```

Fit Multinomial Log-linear Models

Description:

Fits multinomial log-linear models via neural networks.

Usage:

```
multinom(formula, data, weights, subset, na.action,
          contrasts = NULL, Hess = FALSE, summ = 0, censored = FALSE,
          model = FALSE, ...)
```

Arguments:

formula: a formula expression as for regression models, of the form 'response ~ predictors'. The response should be a factor or a matrix with K columns, which will be interpreted as counts for each of K classes. A log-linear model is fitted, with coefficients zero for the first class. An offset can be included: it should be a numeric matrix with K columns if the response is either a matrix with K columns or a factor with K >= 2 classes, or a numeric vector for a response factor with

2 levels. See the documentation of 'formula()' for other details.

data: an optional data frame in which to interpret the variables occurring in 'formula'.

weights: optional case weights in fitting.

subset: expression saying which subset of the rows of the data should be used in the fit. All observations are included by default.

na.action: a function to filter missing data.

contrasts: a list of contrasts to be used for some or all of the factors appearing as variables in the model formula.

Hess: logical for whether the Hessian (the observed/expected information matrix) should be returned.

summ: integer; if non-zero summarize by deleting duplicate rows and adjust weights. Methods 1 and 2 differ in speed (2 uses 'C'); method 3 also combines rows with the same X and different Y, which changes the baseline for the deviance.

censored: If Y is a matrix with 'K' columns, interpret the entries as one for possible classes, zero for impossible classes, rather than as counts.

model: logical. If true, the model frame is saved as component 'model' of the returned object.

...: additional arguments for 'nnet'

For this function the variables are supposed to be roughly 0 to 1. The fit is a little weird if they are scaled or if they are not scaled. Here, I don't scale them.

```
fit <- multinom(Species ~ ., data = iris)
```

```
## # weights: 18 (10 variable)
## initial value 164.791843
## iter 10 value 16.177348
## iter 20 value 7.111438
## iter 30 value 6.182999
## iter 40 value 5.984028
## iter 50 value 5.961278
## iter 60 value 5.954900
## iter 70 value 5.951851
## iter 80 value 5.950343
## iter 90 value 5.949904
## iter 100 value 5.949867
## final value 5.949867
## stopped after 100 iterations
```

```
# summarize the fit
summary(fit)
```

```
## Call:
## multinom(formula = Species ~ ., data = iris)
##
## Coefficients:
##          (Intercept) Sepal.Length Sepal.Width Petal.Length Petal.Width
## versicolor      18.69037      -5.458424      -8.707401       14.24477      -3.097684
## virginica       -23.83628      -7.923634     -15.370769       23.65978       15.135301
##
## Std. Errors:
##          (Intercept) Sepal.Length Sepal.Width Petal.Length Petal.Width
## versicolor      34.97116      89.89215      157.0415       60.19170       45.48852
## virginica       35.76649      89.91153      157.1196       60.46753       45.93406
##
## Residual Deviance: 11.89973
## AIC: 31.89973
```

Now we'll predict for the remaining subjects:

```
pred_log <- predict(fit, newdata = iris[-train,], "probs")
prec_class <- apply(pred_log, 1, which.max)
tbl_2 <- table(prec_class, iris$Species[-train])
prop.table(tbl_2)
```

prec_class/	setosa	versicolor	virginica
1	0.38	0.0	0.00
2	0.00	0.3	0.02
3	0.00	0.0	0.30

```
sum(diag(prop.table(tbl_2)))
```

```
## [1] 0.98
```

```
library(glmnet)
```

Loading required package: Matrix

Loaded glmnet 4.1-8

```
## Here we fit the LASSO model to our data. We'll look at this for the default
## nlambda = 100, and lambda.min.ratio = 0.0001.
`?`(glmnet)
```

fit a GLM with lasso or elasticnet regularization

Description:

Fit a generalized linear model via penalized maximum likelihood. The regularization path is computed for the lasso or elasticnet penalty at a grid of values for the regularization parameter lambda. Can deal with all shapes of data, including very large sparse data matrices. Fits linear, logistic and multinomial, poisson, and Cox regression models.

Usage:

```
glmnet(
```

```

x,
y,
family = c("gaussian", "binomial", "poisson", "multinomial", "cox", "mgaussian"),
weights = NULL,
offset = NULL,
alpha = 1,
nlambda = 100,
lambda.min.ratio = ifelse(nobs < nvars, 0.01, 1e-04),
lambda = NULL,
standardize = TRUE,
intercept = TRUE,
thresh = 1e-07,
dfmax = nvars + 1,
pmax = min(dfmax * 2 + 20, nvars),
exclude = NULL,
penalty.factor = rep(1, nvars),
lower.limits = -Inf,
upper.limits = Inf,
maxit = 1e+05,
type.gaussian = ifelse(nvars < 500, "covariance", "naive"),
type.logistic = c("Newton", "modified.Newton"),
standardize.response = FALSE,
type.multinomial = c("ungrouped", "grouped"),
relax = FALSE,
trace.it = 0,
...
)

```

```
relax.glmnet(fit, x, ..., maxp = n - 3, path = FALSE, check.args = TRUE)
```



```
head(iris)
```


Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.0	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.0	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa

```

lasso_mod <- glmnet(as.matrix(iris[train, 1:4]), iris[train, 5], family = "multinomial",
  alpha = 1)
plot(lasso_mod, xvar = "lambda", label = TRUE)


```





This gives a look at the path of the coefficients. Of course to fit the lasso model we need to select the lambda parameter. Here we use CV to implement this.

```
set.seed(12412)
cvfit = cv.glmnet(as.matrix(iris[train,1:4]), iris[train,5],
                  family="multinomial", alpha=1)
plot(cvfit)
```




```

cvfit$lambda.min

## [1] 0.0007630284

coef(cvfit, s = "lambda.min")

## $setosa
## 5 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) 16.239223
## Sepal.Length .
## Sepal.Width  3.731525
## Petal.Length -2.894364
## Petal.Width  -2.607045
##
## $versicolor
## 5 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept)  9.05159
## Sepal.Length 1.66568
## Sepal.Width  .
## Petal.Length .
## Petal.Width  .
##
## $virginica
## 5 x 1 sparse Matrix of class "dgCMatrix"
##              1
## (Intercept) -25.290813
## Sepal.Length .
## Sepal.Width  -6.277133
## Petal.Length  6.633751
## Petal.Width  17.621958

logit_pred <- predict(cvfit, newx = as.matrix(iris[-train,1:4]), s = "lambda.min")
head(logit_pred)

## , , 1
##
##      setosa versicolor virginica
## 2  22.86028   17.21342 -31.31057
## 4  22.94400   16.71372 -31.27491
## 7  24.09219   16.71372 -32.05923
## 8  24.06345   17.37999 -33.15805
## 9  22.48713   16.38058 -30.68286
## 10 23.20470   17.21342 -33.03710

prec_class_lasso <- apply(logit_pred,1,which.max)
tbl_2 <- table(prec_class_lasso,iris$Species[-train])
prop.table(tbl_2)

```

prec_class_lasso/	setosa	versicolor	virginica
1	0.38	0.0	0.00
2	0.00	0.3	0.02
3	0.00	0.0	0.30

```
sum(diag(prop.table(tbl_2)))
```

```
## [1] 0.98
```