

Neural Networks

ACM

November 8th, 2023

```
## Registered S3 method overwritten by 'printr':  
##   method          from  
##   knit_print.data.frame rmarkdown
```

Our example data set is from the Wisconsin cancer study. We read in the data and remove any rows with missing data. The following code uses the package `mlbench` that contains this data set. We show the top few lines of the data.

```
library(mlbench)  
data("BreastCancer")  
  
#Clean off rows with missing data  
BreastCancer = BreastCancer[which(complete.cases(BreastCancer)==TRUE),]  
dim(BreastCancer)
```

```
## [1] 683 11  
head(BreastCancer)
```

| Id | Cl.thickness | Cell.size | Cell.shape | Marg.adhesion | Epith.c.size | Bare.nuclei | Bl.cromatin | Normal.nuclei | Mitoses | Class |
|----------|--------------|-----------|------------|---------------|--------------|-------------|-------------|---------------|---------|-----------|
| 10000255 | 1 | 1 | 1 | 2 | 1 | 3 | 1 | 1 | 1 | benign |
| 10029455 | 4 | 4 | 5 | 7 | 10 | 3 | 2 | 1 | 1 | benign |
| 10154253 | 1 | 1 | 1 | 2 | 2 | 3 | 1 | 1 | 1 | benign |
| 10162776 | 8 | 8 | 1 | 3 | 4 | 3 | 7 | 1 | 1 | benign |
| 10170234 | 1 | 1 | 3 | 2 | 1 | 3 | 1 | 1 | 1 | benign |
| 10171228 | 10 | 10 | 8 | 7 | 10 | 9 | 7 | 1 | 1 | malignant |

?BreastCancer

```
## Wisconsin Breast Cancer Database  
##  
## Description:  
##  
##   The objective is to identify each of a number of benign or  
##   malignant classes. Samples arrive periodically as Dr. Wolberg  
##   reports his clinical cases. The database therefore reflects this  
##   chronological grouping of the data. This grouping information  
##   appears immediately below, having been removed from the data  
##   itself. Each variable except for the first was converted into 11  
##   primitive numerical attributes with values ranging from 0 through  
##   10. There are 16 missing attribute values. See cited below for  
##   more details.  
##  
## Format:
```

```
##
##      A data frame with 699 observations on 11 variables, one being a
##      character variable, 9 being ordered or nominal, and 1 target
##      class.
```

```
##      [,1] Id      Sample code number
##      [,2] Cl.thickness Clump Thickness
##      [,3] Cell.size  Uniformity of Cell Size
##      [,4] Cell.shape Uniformity of Cell Shape
##      [,5] Marg.adhesion Marginal Adhesion
##      [,6] Epith.c.size Single Epithelial Cell Size
##      [,7] Bare.nuclei Bare Nuclei
##      [,8] Bl.cromatin Bland Chromatin
##      [,9] Normal.nucleoli Normal Nucleoli
##      [,10] Mitoses    Mitoses
##      [,11] Class      Class
```

```
str(BreastCancer)
```

```
## 'data.frame':  683 obs. of  11 variables:
## $ Id      : chr  "1000025" "1002945" "1015425" "1016277" ...
## $ Cl.thickness : Ord.factor w/ 10 levels "1"<"2"<"3"<"4"<..: 5 5 3 6 4 8 1 2 2 4 ...
## $ Cell.size  : Ord.factor w/ 10 levels "1"<"2"<"3"<"4"<..: 1 4 1 8 1 10 1 1 1 2 ...
## $ Cell.shape : Ord.factor w/ 10 levels "1"<"2"<"3"<"4"<..: 1 4 1 8 1 10 1 2 1 1 ...
## $ Marg.adhesion : Ord.factor w/ 10 levels "1"<"2"<"3"<"4"<..: 1 5 1 1 3 8 1 1 1 1 ...
## $ Epith.c.size : Ord.factor w/ 10 levels "1"<"2"<"3"<"4"<..: 2 7 2 3 2 7 2 2 2 2 ...
## $ Bare.nuclei  : Factor w/ 10 levels "1","2","3","4",...: 1 10 2 4 1 10 10 1 1 1 ...
## $ Bl.cromatin  : Factor w/ 10 levels "1","2","3","4",...: 3 3 3 3 3 9 3 3 1 2 ...
## $ Normal.nucleoli: Factor w/ 10 levels "1","2","3","4",...: 1 2 1 7 1 7 1 1 1 1 ...
## $ Mitoses      : Factor w/ 9 levels "1","2","3","4",...: 1 1 1 1 1 1 1 1 5 1 ...
## $ Class        : Factor w/ 2 levels "benign","malignant": 1 1 1 1 1 2 1 1 1 1 ...
```

Since `neuralnet` only deals with quantitative variables, you can convert all the qualitative variables (factors) to binary (“dummy”), with the `model.matrix` function.

```
y = as.matrix(BreastCancer[,11])
y[which(y=="benign")] = 0
y[which(y=="malignant")] = 1
y = as.numeric(y)
x = model.matrix(~0 + Cl.thickness + Cell.size + Cell.shape +
                Marg.adhesion + Epith.c.size + Bare.nuclei + Bl.cromatin +
                Normal.nucleoli + Mitoses, data = BreastCancer)
dim(x)
```

```
## [1] 683 81
```

```
df <- data.frame(y,x)
```

Let’s split the data into training and validation:

```
set.seed(92739847)
ids <- sample(1:dim(df)[1],size=dim(df)[1])
df_jumbled <- df[ids,]
train <- df_jumbled[1:450,]
test <- df_jumbled[451:nrow(df_jumbled),]
```

Fitting with neuralnet

In the first example, we will create a simple neural network with minimum effort using the **neuralnet** package. The **neuralnet** package is outdated, but it is still popular among the R community.

The **neuralnet** function is simple. It doesn't provide us the freedom to create fully customizable model architecture. In our case, we are providing it with a machine-learning formula and data, just like GLM. The formula consists of target variables and features.

```
library("neuralnet")
?neuralnet
```

Training of neural networks

Description:

Train neural networks using backpropagation, resilient backpropagation (RPROP) with (Riedmiller, 1994) or without weight backtracking (Riedmiller and Braun, 1993) or the modified globally convergent version (GRPROP) by Anastasiadis et al. (2005). The function allows flexible settings through custom-choice of error and activation function. Furthermore, the calculation of generalized weights (Intrator O. and Intrator N., 1993) is implemented.

Usage:

```
neuralnet(formula, data, hidden = 1, threshold = 0.01,
  stepmax = 1e+05, rep = 1, startweights = NULL,
  learningrate.limit = NULL, learningrate.factor = list(minus = 0.5,
  plus = 1.2), learningrate = NULL, lifesign = "none",
  lifesign.step = 1000, algorithm = "rprop+", err.fct = "sse",
  act.fct = "logistic", linear.output = TRUE, exclude = NULL,
  constant.weights = NULL, likelihood = FALSE)
```

Arguments:

formula: a symbolic description of the model to be fitted.

data: a data frame containing the variables specified in 'formula'.

hidden: a vector of integers specifying the number of hidden neurons (vertices) in each layer.

threshold: a numeric value specifying the threshold for the partial derivatives of the error function as stopping criteria.

stepmax: the maximum steps for the training of the neural network. Reaching this maximum leads to a stop of the neural network's training process.

rep: the number of repetitions for the neural network's training.

startweights: a vector containing starting values for the weights. Set to 'NULL' for random initialization.

`learningrate.limit`: a vector or a list containing the lowest and highest limit for the learning rate. Used only for RPROP and GRPROP.

`learningrate.factor`: a vector or a list containing the multiplication factors for the upper and lower learning rate. Used only for RPROP and GRPROP.

`learningrate`: a numeric value specifying the learning rate used by traditional backpropagation. Used only for traditional backpropagation.

`lifesign`: a string specifying how much the function will print during the calculation of the neural network. 'none', 'minimal' or 'full'.

`lifesign.step`: an integer specifying the stepsize to print the minimal threshold in full lifesign mode.

`algorithm`: a string containing the algorithm type to calculate the neural network. The following types are possible: 'backprop', 'rprop+', 'rprop-', 'sag', or 'slr'. 'backprop' refers to backpropagation, 'rprop+' and 'rprop-' refer to the resilient backpropagation with and without weight backtracking, while 'sag' and 'slr' induce the usage of the modified globally convergent algorithm (grprop). See Details for more information.

`err.fct`: a differentiable function that is used for the calculation of the error. Alternatively, the strings 'sse' and 'ce' which stand for the sum of squared errors and the cross-entropy can be used.

`act.fct`: a differentiable function that is used for smoothing the result of the cross product of the covariate or neurons and the weights. Additionally the strings, 'logistic' and 'tanh' are possible for the logistic function and tangent hyperbolicus.

`linear.output`: logical. If `act.fct` should not be applied to the output neurons set linear output to TRUE, otherwise to FALSE.

`exclude`: a vector or a matrix specifying the weights, that are excluded from the calculation. If given as a vector, the exact positions of the weights must be known. A matrix with n-rows and 3 columns will exclude n weights, where the first column stands for the layer, the second column for the input neuron and the third column for the output neuron of the weight.

`constant.weights`: a vector specifying the values of the weights that are excluded from the training process and treated as fix.

likelihood: logical. If the error function is equal to the negative log-likelihood function, the information criteria AIC and BIC will be calculated. Furthermore the usage of confidence.interval is meaningful.

Details:

The globally convergent algorithm is based on the resilient backpropagation without weight backtracking and additionally modifies one learning rate, either the learningrate associated with the smallest absolute gradient (sag) or the smallest learningrate (slr) itself. The learning rates in the grprop algorithm are limited to the boundaries defined in learningrate.limit.

Now we'll fit a neural network. How many layers?

- There's a lot of discussion on this. See <http://www.faqs.org/faqs/ai-faq/neural-nets/part1/preamble.html> for some.
- One thing that is agreed upon is that the situations in which performance improves with a second (or third, etc.) hidden layer are very small.
- **One hidden layer is sufficient for the large majority of problems.**

How about the size of the hidden layer(s):

- A common rule of thumb is that the optimal size of the hidden layer is usually between the size of the input and size of the output layers.
- One way to choose would be to use the mean.

In this example we'll use 81 input layers and 1 output layer, so we'll choose 40.

```
Breast_NN <- neuralnet(y ~ ., data = train,
                      hidden = 40, lifesign = "minimal", algorithm="backprop",
                      linear.output = FALSE, err.fct = 'ce',
                      learningrate=0.1)
```

```
## hidden: 40    thresh: 0.01    rep: 1/1    steps:    2282 error: 0.04201    time: 9.63 secs
```

```
# Look at the error and activation function
```

```
Breast_NN$err.fct
```

```
## function (x, y)
## {
##     -(y * log(x) + (1 - y) * log(1 - x))
## }
## <bytecode: 0x111558fc0>
## <environment: 0x11155c060>
## attr("type")
## [1] "ce"
```

```
Breast_NN$act.fct
```

```
## function (x)
## {
##     1/(1 + exp(-x))
## }
## <bytecode: 0x11154e4b0>
## <environment: 0x111551ac8>
## attr("type")
```

```
## [1] "logistic"
```

```
# Here are the parameters:
```

```
dim(Breast_NN$result.matrix)
```

```
## [1] 3324    1
```

```
3321/81
```

```
## [1] 41
```

```
head(data.frame(Breast_NN$result.matrix))
```

| | Breast_NN.result.matrix |
|---------------------------|-------------------------|
| error | 0.0420063 |
| reached.threshold | 0.0099975 |
| steps | 2282.0000000 |
| Intercept.to.1layhid1 | -7.3351334 |
| Cl.thickness1.to.1layhid1 | 1.9221154 |
| Cl.thickness2.to.1layhid1 | 1.6221521 |

```
#plot(Breast_NN, rep="best")
```



```
valid_pred <- predict(Breast_NN, newdata = test)
table(valid_pred>0.5, test$y)
```

| / | 0 | 1 |
|-------|-----|----|
| FALSE | 152 | 9 |
| TRUE | 1 | 71 |

```
# Accuracy
sum(diag(table(valid_pred>0.5, test$y)))/length(test$y)
```

```
## [1] 0.9570815
```

Fitting with TensorFlow

TensorFlow is an open source deep neural net framework, based on a graphical model. It is more than just a neural net platform, and supports numerical computing based on data flow graphs. Data may be represented in n -dimensional structures like vectors and matrices, or higher-dimensional tensors. Because these mathematical objects are folded into a data flow graph for computation, the moniker for this software library is an obvious one.

The computations for deep learning nets involve tensor computations, which are known to be implemented more efficiently on GPUs than CPUs. Therefore like other deep learning libraries, **TensorFlow** may be implemented on CPUs and GPUs. The generality and speed of the **TensorFlow** software, ease of installation, its documentation and examples, and runnability on multiple platforms has made **TensorFlow** the most popular deep learning toolkit today. (Opinions on this may, of course, differ.) There is a wealth of tutorial material on TF of very high quality that you may refer to here: <https://www.tensorflow.org/tutorials/>.

Rather than run **TensorFlow** natively, it is often easier to use it through an easy to use interface program. One of the most popular high-level APIs is **Keras** See: <https://keras.io/>. The site contains several examples, which make it easy to get up and running. Though originally written in Python, **Keras** has been extended to R via the **keras** package.

Keras is a high-level neural networks API developed with a focus on enabling fast experimentation. Being able to go from idea to result with the least possible delay is key to doing good research. Keras has the following key features:

- Allows the same code to run on CPU or on GPU, seamlessly.
- User-friendly API which makes it easy to quickly prototype deep learning models.
- Built-in support for convolutional networks (for computer vision), recurrent networks (for sequence processing), and any combination of both.
- Supports arbitrary network architectures: multi-input or multi-output models, layer sharing, model sharing, etc. This means that Keras is appropriate for building essentially any deep learning model, from a memory network to a neural Turing machine.

```
library(tensorflow)
library(keras)
library(reticulate)
```

To use these functions were going to convert the outcome (benign, malignant) in the data set to “one-hot encoding” using the `to_categorical` function. What is one-hot encoding replaces a single vector of 1s (malignant) and 0s (benign) to a bivariate dependent variable as a two-column matrix, where the first column contains a 1 if the cell is benign and the second column a 1 if the cell is malignant, for each row of the data set. This is because TensorFlow/Keras requires this format as input, to facilitate tensor calculations.


```

y_train = as.matrix(train[,1])
y_test = as.matrix(test[,1])
x_train = as.matrix(train[,-1])
x_test = as.matrix(test[,-1])

Y_train = to_categorical(y_train,2)
Y_test = to_categorical(y_test,2)

```

TensorFlow is structured in a manner where one describes the neural net first, layer by layer. There is not a single function that is called, which generates the deep learning net, and runs the model. Here, we first describe for each layer in the neural net, the number of nodes, the type of activation function, and any other hyperparameters needed in the model fitting stage, such as the extent of dropout for example.

In the output layer we also state the nature of the activation function, such as sigmoid or softmax. And then, we specify a compile function which describes the loss function that will be minimized, along with the minimization algorithm. At this point in the program specification, the model is not actually run. See below for the code block that builds up the deep learning network.

The core data structure of Keras is a model, a way to organize layers. The simplest type of model is the Sequential model, a linear stack of layers.

We begin by creating a sequential model and then adding layers using the pipe (%>%) operator:

```
?keras_model_sequential
```

Keras Model composed of a linear stack of layers

Description:

Keras Model composed of a linear stack of layers

Usage:

```
keras_model_sequential(layers = NULL, name = NULL, ...)
```

Arguments:

layers: List of layers to add to the model

name: Name of model

...: Arguments passed on to 'sequential_model_input_layer'

'input_shape' an integer vector of dimensions (not including the batch axis), or a 'tf\$TensorShape' instance (also not including the batch axis).

'batch_size' Optional input batch size (integer or NULL).

'dtype' Optional datatype of the input. When not provided, the Keras default float type will be used.

'input_tensor' Optional tensor to use as layer input. If set, the layer will use the 'tf\$TypeSpec' of this tensor rather than creating a new placeholder tensor.

'sparse' Boolean, whether the placeholder created is meant to be sparse. Default to 'FALSE'.

'ragged' Boolean, whether the placeholder created is meant to be ragged. In this case, values of 'NULL' in the 'shape' argument represent ragged dimensions. For more information about 'RaggedTensors', see this guide. Default to 'FALSE'.

'type_spec' A 'tf\$TypeSpec' object to create Input from. This 'tf\$TypeSpec' represents the entire batch. When provided, all other args except name must be 'NULL'.

'input_layer_name,name' Optional name of the input layer (string).

```
model <- keras_model_sequential()
n <- nrow(x_train)
model %>%
  layer_dense(units = n/2, activation = 'relu', input_shape = c(ncol(x_train))) %>%
  # set rate: the fraction of the input units to drop.
  layer_dropout(rate = 0.25) %>%
  layer_dense(units = n/10, activation = 'relu') %>%
  layer_dropout(rate = 0.2) %>%
  layer_dense(units = 2, activation = 'softmax')
```

The `input_shape` argument to the first layer specifies the shape of the input data (a length `ncol(x_train)` numeric vector). The final layer outputs a length 2 numeric vector (probabilities for each class) using a softmax activation function.

```
summary(model)
```

```
## Model: "sequential"
## -----
## Layer (type)                Output Shape          Param #
## -----
## dense_2 (Dense)              (None, 225)           18450
## dropout_1 (Dropout)          (None, 225)           0
## dense_1 (Dense)              (None, 45)            10170
## dropout (Dropout)            (None, 45)            0
## dense (Dense)                (None, 2)             92
## -----
## Total params: 28712 (112.16 KB)
## Trainable params: 28712 (112.16 KB)
## Non-trainable params: 0 (0.00 Byte)
## -----
```

Next, compile the model with appropriate loss function, optimizer, and metrics:

```
?compile.keras.engine.training.Model
```

Configure a Keras model for training

Description:

Configure a Keras model for training

Usage:

```
## S3 method for class 'keras.engine.training.Model'
compile(
  object,
  optimizer = NULL,
  loss = NULL,
  metrics = NULL,
  loss_weights = NULL,
  weighted_metrics = NULL,
  run_eagerly = NULL,
  steps_per_execution = NULL,
  ...,
  target_tensors = NULL,
  sample_weight_mode = NULL
)
```

Arguments:

object: Model object to compile.

optimizer: String (name of optimizer) or optimizer instance. For most models, this defaults to "rmsprop"

loss: String (name of objective function), objective function or a 'keras\$losses\$Loss' subclass instance. An objective function is any callable with the signature 'loss = fn(y_true, y_pred)', where y_true = ground truth values with shape = [batch_size, d0, .. dN], except sparse loss functions such as sparse categorical crossentropy where shape = [batch_size, d0, .. dN-1]. y_pred = predicted values with shape = [batch_size, d0, .. dN]. It returns a weighted loss float tensor. If a custom 'Loss' instance is used and reduction is set to 'NULL', return value has the shape [batch_size, d0, .. dN-1] i.e. per-sample or per-timestep loss values; otherwise, it is a scalar. If the model has multiple outputs, you can use a different loss on each output by passing a dictionary or a list of losses. The loss value that will be minimized by the model will then be the sum of all individual losses, unless 'loss_weights' is specified.

metrics: List of metrics to be evaluated by the model during training and testing. Each of this can be a string (name of a built-in function), function or a 'keras\$metrics\$Metric' class instance. See '?tf\$keras\$metrics'. Typically you will use 'metrics=list('accuracy')'. A function is any callable with the signature 'result = fn(y_true, y_pred)'. To specify different metrics for different outputs of a multi-output model, you could also pass a dictionary, such as 'metrics=list(output_a = 'accuracy', output_b = c('accuracy', 'mse'))'. You can also pass a list to specify a metric or a list of metrics for each output, such as 'metrics=list(list('accuracy'), list('accuracy', 'mse'))' or 'metrics=list('accuracy', c('accuracy', 'mse'))'. When you

pass the strings 'accuracy' or 'acc', this is converted to one of 'tf.keras.metrics.BinaryAccuracy', 'tf.keras.metrics.CategoricalAccuracy', 'tf.keras.metrics.SparseCategoricalAccuracy' based on the loss function used and the model output shape. A similar conversion is done for the strings 'crossentropy' and 'ce'.

loss_weights: Optional list, dictionary, or named vector specifying scalar numeric coefficients to weight the loss contributions of different model outputs. The loss value that will be minimized by the model will then be the `_weighted sum_` of all individual losses, weighted by the 'loss_weights' coefficients. If a list, it is expected to have a 1:1 mapping to the model's outputs. If a dict, it is expected to map output names (strings) to scalar coefficients.

weighted_metrics: List of metrics to be evaluated and weighted by 'sample_weight' or 'class_weight' during training and testing.

run_eagerly: Bool. Defaults to 'FALSE'. If 'TRUE', this Model's logic will not be wrapped in a 'tf.function'. Recommended to leave this as 'NULL' unless your Model cannot be run inside a 'tf.function'. 'run_eagerly=True' is not supported when using 'tf.distribute.experimental.ParameterServerStrategy'. If the model's logic uses tensors in R control flow expressions like 'if' and 'for', the model is still traceable with 'tf.function', but you will have to enter a 'tfautograph::autograph({})' directly.

steps_per_execution: Int. Defaults to 1. The number of batches to run during each 'tf.function' call. Running multiple batches inside a single 'tf.function' call can greatly improve performance on TPUs or small models with a large Python/R overhead. At most, one full epoch will be run each execution. If a number larger than the size of the epoch is passed, the execution will be truncated to the size of the epoch. Note that if 'steps_per_execution' is set to 'N', 'Callback.on_batch_begin' and 'Callback.on_batch_end' methods will only be called every 'N' batches (i.e. before/after each 'tf.function' execution).

...: Arguments supported for backwards compatibility only.

target_tensors: By default, Keras will create a placeholder for the model's target, which will be fed with the target data during training. If instead you would like to use your own target tensor (in turn, Keras will not expect external data for these targets at training time), you can specify them via the 'target_tensors' argument. It should be a single tensor (for a single-output sequential model).

sample_weight_mode: If you need to do timestep-wise sample weighting

(2D weights), set this to "temporal". 'NULL' defaults to sample-wise weights (1D). If the model has multiple outputs, you can use a different 'sample_weight_mode' on each output by passing a list of modes.

```
model %>% compile(  
  loss = 'categorical_crossentropy',  
  optimizer = optimizer_rmsprop(),  
  metrics = c('accuracy')  
)
```

From this I got the following message:

WARNING:absl:At this time, the v2.11+ optimizer `tf.keras.optimizers.RMSprop` runs slowly on M1/M2 Macs, please use the legacy Keras optimizer instead, located at `tf.keras.optimizers.legacy.RMSprop`.

This error apparently happens when trying to use ADAM (on python or R). The following does work:

```
model %>% compile(  
  loss = 'categorical_crossentropy',  
  metrics = c('accuracy')  
)
```

Training and Evaluation

Use the `fit()` function to train the model for 30 epochs using batches of 90 subjects:

```
?fit.keras.engine.training.Model
```

Train a Keras model

Description:

Trains the model for a fixed number of epochs (iterations on a dataset).

Usage:

```
## S3 method for class 'keras.engine.training.Model'  
fit(  
  object,  
  x = NULL,  
  y = NULL,  
  batch_size = NULL,  
  epochs = 10,  
  verbose = getOption("keras.fit_verbose", default = "auto"),  
  callbacks = NULL,  
  view_metrics = getOption("keras.view_metrics", default = "auto"),  
  validation_split = 0,  
  validation_data = NULL,  
  shuffle = TRUE,  
  class_weight = NULL,  
  sample_weight = NULL,  
  initial_epoch = 0,  
  steps_per_epoch = NULL,  
  validation_steps = NULL,  
  ...  
)
```

Arguments:

- `object`: Model to train.
- `x`: Vector, matrix, or array of training data (or list if the model has multiple inputs). If all inputs in the model are named, you can also pass a list mapping input names to data. 'x' can be 'NULL' (default) if feeding from framework-native tensors (e.g. TensorFlow data tensors). You can also pass a 'tfdataset' or a generator returning a list with (inputs, targets) or (inputs, targets, sample_weights).
- `y`: Vector, matrix, or array of target (label) data (or list if the model has multiple outputs). If all outputs in the model are named, you can also pass a list mapping output names to data. 'y' can be 'NULL' (default) if feeding from framework-native tensors (e.g. TensorFlow data tensors).
- `batch_size`: Integer or 'NULL'. Number of samples per gradient update. If unspecified, 'batch_size' will default to 32.
- `epochs`: Number of epochs to train the model. Note that in conjunction with 'initial_epoch', 'epochs' is to be understood as "final epoch". The model is not trained for a number of iterations given by 'epochs', but merely until the epoch of index 'epochs' is reached.
- `verbose`: Verbosity mode (0 = silent, 1 = progress bar, 2 = one line per epoch). Defaults to 1 in most contexts, 2 if in knitr render or running on a distributed training server.
- `callbacks`: List of callbacks to be called during training.
- `view_metrics`: View realtime plot of training metrics (by epoch). The default ('auto') will display the plot when running within RStudio, 'metrics' were specified during model 'compile()', 'epochs > 1' and 'verbose > 0'. Use the global 'keras.view_metrics' option to establish a different default.
- `validation_split`: Float between 0 and 1. Fraction of the training data to be used as validation data. The model will set apart this fraction of the training data, will not train on it, and will evaluate the loss and any model metrics on this data at the end of each epoch. The validation data is selected from the last samples in the 'x' and 'y' data provided, before shuffling.
- `validation_data`: Data on which to evaluate the loss and any model metrics at the end of each epoch. The model will not be trained on this data. This could be a list (x_val, y_val) or a list (x_val, y_val, val_sample_weights). 'validation_data' will override 'validation_split'.

shuffle: shuffle: Logical (whether to shuffle the training data before each epoch) or string (for "batch"). "batch" is a special option for dealing with the limitations of HDF5 data; it shuffles in batch-sized chunks. Has no effect when 'steps_per_epoch' is not 'NULL'.

class_weight: Optional named list mapping indices (integers) to a weight (float) value, used for weighting the loss function (during training only). This can be useful to tell the model to "pay more attention" to samples from an under-represented class.

sample_weight: Optional array of the same length as x, containing weights to apply to the model's loss for each sample. In the case of temporal data, you can pass a 2D array with shape (samples, sequence_length), to apply a different weight to every timestep of every sample. In this case you should make sure to specify 'sample_weight_mode="temporal"' in 'compile()'.

initial_epoch: Integer, Epoch at which to start training (useful for resuming a previous training run).

steps_per_epoch: Total number of steps (batches of samples) before declaring one epoch finished and starting the next epoch. When training with input tensors such as TensorFlow data tensors, the default 'NULL' is equal to the number of samples in your dataset divided by the batch size, or 1 if that cannot be determined.

validation_steps: Only relevant if 'steps_per_epoch' is specified. Total number of steps (batches of samples) to validate before stopping.

...: Unused

```
history <- model %>% fit(
  x_train, Y_train,
  epochs = 30, batch_size = 90,
  validation_split = 0.2
)
```

Epoch 1/30

4/4 - 1s - loss: 0.5306 - accuracy: 0.7694 - val_loss: 0.3525 - val_accuracy: 0.9556 - 715ms/epoch - 17s

Epoch 2/30

4/4 - 0s - loss: 0.3185 - accuracy: 0.9417 - val_loss: 0.2315 - val_accuracy: 0.9667 - 71ms/epoch - 18ms

Epoch 3/30

4/4 - 0s - loss: 0.2120 - accuracy: 0.9611 - val_loss: 0.1731 - val_accuracy: 0.9556 - 40ms/epoch - 10ms

Epoch 4/30

4/4 - 0s - loss: 0.1671 - accuracy: 0.9639 - val_loss: 0.1442 - val_accuracy: 0.9556 - 45ms/epoch - 11ms

Epoch 5/30

4/4 - 0s - loss: 0.1239 - accuracy: 0.9694 - val_loss: 0.1293 - val_accuracy: 0.9556 - 60ms/epoch - 15ms

Epoch 6/30

4/4 - 0s - loss: 0.1081 - accuracy: 0.9778 - val_loss: 0.1183 - val_accuracy: 0.9556 - 48ms/epoch - 12ms

Epoch 7/30


```

4/4 - 0s - loss: 0.0973 - accuracy: 0.9778 - val_loss: 0.1143 - val_accuracy: 0.9556 - 45ms/epoch - 11ms
Epoch 8/30
4/4 - 0s - loss: 0.0839 - accuracy: 0.9778 - val_loss: 0.1073 - val_accuracy: 0.9667 - 48ms/epoch - 12ms
Epoch 9/30
4/4 - 0s - loss: 0.0629 - accuracy: 0.9806 - val_loss: 0.1120 - val_accuracy: 0.9667 - 40ms/epoch - 10ms
Epoch 10/30
4/4 - 0s - loss: 0.0593 - accuracy: 0.9833 - val_loss: 0.1139 - val_accuracy: 0.9667 - 40ms/epoch - 10ms
Epoch 11/30
4/4 - 0s - loss: 0.0576 - accuracy: 0.9806 - val_loss: 0.1208 - val_accuracy: 0.9667 - 32ms/epoch - 8ms
Epoch 12/30
4/4 - 0s - loss: 0.0463 - accuracy: 0.9806 - val_loss: 0.1184 - val_accuracy: 0.9667 - 36ms/epoch - 9ms
Epoch 13/30
4/4 - 0s - loss: 0.0350 - accuracy: 0.9917 - val_loss: 0.1261 - val_accuracy: 0.9667 - 40ms/epoch - 10ms
Epoch 14/30
4/4 - 0s - loss: 0.0353 - accuracy: 0.9889 - val_loss: 0.1305 - val_accuracy: 0.9667 - 34ms/epoch - 8ms
Epoch 15/30
4/4 - 0s - loss: 0.0259 - accuracy: 0.9917 - val_loss: 0.1333 - val_accuracy: 0.9667 - 25ms/epoch - 6ms
Epoch 16/30
4/4 - 0s - loss: 0.0271 - accuracy: 0.9944 - val_loss: 0.1347 - val_accuracy: 0.9667 - 39ms/epoch - 10ms
Epoch 17/30
4/4 - 0s - loss: 0.0219 - accuracy: 0.9944 - val_loss: 0.1516 - val_accuracy: 0.9667 - 39ms/epoch - 10ms
Epoch 18/30
4/4 - 0s - loss: 0.0197 - accuracy: 0.9972 - val_loss: 0.1589 - val_accuracy: 0.9667 - 30ms/epoch - 7ms
Epoch 19/30
4/4 - 0s - loss: 0.0168 - accuracy: 0.9944 - val_loss: 0.1724 - val_accuracy: 0.9667 - 31ms/epoch - 8ms
Epoch 20/30
4/4 - 0s - loss: 0.0120 - accuracy: 0.9972 - val_loss: 0.1728 - val_accuracy: 0.9667 - 31ms/epoch - 8ms
Epoch 21/30
4/4 - 0s - loss: 0.0115 - accuracy: 0.9972 - val_loss: 0.1761 - val_accuracy: 0.9667 - 64ms/epoch - 16ms
Epoch 22/30
4/4 - 0s - loss: 0.0138 - accuracy: 0.9944 - val_loss: 0.1841 - val_accuracy: 0.9667 - 59ms/epoch - 15ms
Epoch 23/30
4/4 - 0s - loss: 0.0092 - accuracy: 0.9944 - val_loss: 0.2183 - val_accuracy: 0.9667 - 43ms/epoch - 11ms
Epoch 24/30
4/4 - 0s - loss: 0.0107 - accuracy: 0.9944 - val_loss: 0.2153 - val_accuracy: 0.9556 - 41ms/epoch - 10ms
Epoch 25/30
4/4 - 0s - loss: 0.0044 - accuracy: 1.0000 - val_loss: 0.2127 - val_accuracy: 0.9667 - 37ms/epoch - 9ms
Epoch 26/30
4/4 - 0s - loss: 0.0038 - accuracy: 1.0000 - val_loss: 0.2225 - val_accuracy: 0.9667 - 44ms/epoch - 11ms
Epoch 27/30
4/4 - 0s - loss: 0.0029 - accuracy: 1.0000 - val_loss: 0.2202 - val_accuracy: 0.9667 - 35ms/epoch - 9ms
Epoch 28/30
4/4 - 0s - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.2172 - val_accuracy: 0.9667 - 87ms/epoch - 22ms
Epoch 29/30
4/4 - 0s - loss: 0.0041 - accuracy: 0.9972 - val_loss: 0.2228 - val_accuracy: 0.9556 - 33ms/epoch - 8ms
Epoch 30/30
4/4 - 0s - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.2313 - val_accuracy: 0.9667 - 25ms/epoch - 6ms

```

The `history` object returned by `fit()` includes loss and accuracy metrics which we can plot:

```
plot(history)
```

Evaluate the model's performance on the test data:

```
model %>% evaluate(x_test, Y_test)
```

```
## 8/8 - 0s - loss: 0.0582 - accuracy: 0.9785 - 108ms/epoch - 14ms/step
```

```
##      loss  accuracy
## 0.05818005 0.97854078
```

Generate predictions on new data:

```
model %>% predict(x_test) %>% k_argmax()
```

```
## 8/8 - 0s - 69ms/epoch - 9ms/step
```

```
## tf.Tensor(
## [1 1 0 1 1 0 1 0 1 0 0 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 1 0 0 0 1
## 1 0 0 0 0 0 0 0 0 0 0 1 0 0 1 0 1 1 0 0 0 0 0 1 0 1 0 0 0 0 0 1 0 0 0 1 0 1 1
## 1 1 0 0 0 1 1 0 0 0 1 0 0 1 0 1 0 0 1 1 0 1 1 0 0 0 1 0 0 0 0 0 1 1 0 1 0
## 0 1 0 0 0 0 0 1 1 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 0 1 0 0 0 1 0 0 0 0 0 0 0
## 1 0 0 1 0 1 1 1 0 0 0 0 0 0 0 1 0 0 0 0 1 0 1 1 0 0 1 0 0 0 1 0 0 0 0 0 0
## 1 0 1 1 0 1 1 0 0 1 1 1 0 0 0 0 0 0 1 1 1 1 1 1 1 0 0 1 0 1 1 0 1 0 1 0 0
## 0 0 0 0 0 1 0 1 0 0 0], shape=(233), dtype=int64)
```

Example with image data

Let's see an example with the mnist image data:

```
mnist <- dataset_mnist()
x_train <- mnist$train$x
```

```
y_train <- mnist$train$y
x_test <- mnist$test$x
y_test <- mnist$test$y
dim(x_train)
```

```
## [1] 60000    28    28
```

```
dim(x_test)
```

```
## [1] 10000    28    28
```

```
# reshape
x_train <- array_reshape(x_train, c(nrow(x_train), 784))
x_test  <- array_reshape(x_test,  c(nrow(x_test),  784))
```

Now we normalize the values, which are pixel intensities ranging in (0,255) and create the outcome.

```
x_train <- x_train / 255
x_test  <- x_test  / 255
y_train <- to_categorical(y_train, 10)
y_test  <- to_categorical(y_test,  10)
```

Now create the model in Keras.

```
model <- keras_model_sequential()
model %>%
  layer_dense(units = 256, activation = 'relu', input_shape = c(784)) %>%
  layer_dropout(rate = 0.4) %>%
  layer_dense(units = 128, activation = 'relu') %>%
  layer_dropout(rate = 0.3) %>%
  layer_dense(units = 10, activation = 'softmax')
summary(model)
```

```
## Model: "sequential_1"
```

```
## -----
## Layer (type)                Output Shape                Param #
## =====
## dense_5 (Dense)              (None, 256)                 200960
## dropout_3 (Dropout)          (None, 256)                 0
## dense_4 (Dense)              (None, 128)                 32896
## dropout_2 (Dropout)          (None, 128)                 0
## dense_3 (Dense)              (None, 10)                  1290
## =====
```

```
## Total params: 235146 (918.54 KB)
## Trainable params: 235146 (918.54 KB)
## Non-trainable params: 0 (0.00 Byte)
## -----
```

```
model %>% compile(
  loss = 'categorical_crossentropy',
  metrics = c('accuracy')
)
```

Now run the model to get a fitted deep learning network.


```
history <- model %>% fit(
  x_train, y_train,
  epochs = 30, batch_size = 128,
```

```
validation_split = 0.2
)
```

```
## Epoch 1/30
## 375/375 - 3s - loss: 0.4241 - accuracy: 0.8738 - val_loss: 0.1553 - val_accuracy: 0.9547 - 3s/epoch
## Epoch 2/30
## 375/375 - 2s - loss: 0.2035 - accuracy: 0.9400 - val_loss: 0.1229 - val_accuracy: 0.9655 - 2s/epoch
## Epoch 3/30
## 375/375 - 3s - loss: 0.1556 - accuracy: 0.9545 - val_loss: 0.1043 - val_accuracy: 0.9703 - 3s/epoch
## Epoch 4/30
## 375/375 - 3s - loss: 0.1321 - accuracy: 0.9608 - val_loss: 0.0936 - val_accuracy: 0.9733 - 3s/epoch
## Epoch 5/30
## 375/375 - 2s - loss: 0.1137 - accuracy: 0.9674 - val_loss: 0.0902 - val_accuracy: 0.9742 - 2s/epoch
## Epoch 6/30
## 375/375 - 2s - loss: 0.1032 - accuracy: 0.9695 - val_loss: 0.0857 - val_accuracy: 0.9772 - 2s/epoch
## Epoch 7/30
## 375/375 - 2s - loss: 0.0955 - accuracy: 0.9715 - val_loss: 0.0890 - val_accuracy: 0.9762 - 2s/epoch
## Epoch 8/30
## 375/375 - 2s - loss: 0.0869 - accuracy: 0.9741 - val_loss: 0.0920 - val_accuracy: 0.9770 - 2s/epoch
## Epoch 9/30
## 375/375 - 2s - loss: 0.0835 - accuracy: 0.9754 - val_loss: 0.0911 - val_accuracy: 0.9779 - 2s/epoch
## Epoch 10/30
## 375/375 - 2s - loss: 0.0801 - accuracy: 0.9767 - val_loss: 0.0907 - val_accuracy: 0.9782 - 2s/epoch
## Epoch 11/30
## 375/375 - 2s - loss: 0.0731 - accuracy: 0.9788 - val_loss: 0.0971 - val_accuracy: 0.9791 - 2s/epoch
## Epoch 12/30
## 375/375 - 2s - loss: 0.0745 - accuracy: 0.9796 - val_loss: 0.0918 - val_accuracy: 0.9784 - 2s/epoch
## Epoch 13/30
## 375/375 - 2s - loss: 0.0718 - accuracy: 0.9799 - val_loss: 0.0900 - val_accuracy: 0.9797 - 2s/epoch
## Epoch 14/30
## 375/375 - 2s - loss: 0.0677 - accuracy: 0.9809 - val_loss: 0.0963 - val_accuracy: 0.9794 - 2s/epoch
## Epoch 15/30
## 375/375 - 2s - loss: 0.0634 - accuracy: 0.9816 - val_loss: 0.0973 - val_accuracy: 0.9804 - 2s/epoch
## Epoch 16/30
## 375/375 - 2s - loss: 0.0646 - accuracy: 0.9814 - val_loss: 0.1001 - val_accuracy: 0.9793 - 2s/epoch
## Epoch 17/30
## 375/375 - 2s - loss: 0.0629 - accuracy: 0.9826 - val_loss: 0.0959 - val_accuracy: 0.9801 - 2s/epoch
## Epoch 18/30
## 375/375 - 2s - loss: 0.0607 - accuracy: 0.9826 - val_loss: 0.1027 - val_accuracy: 0.9805 - 2s/epoch
## Epoch 19/30
## 375/375 - 2s - loss: 0.0580 - accuracy: 0.9840 - val_loss: 0.1050 - val_accuracy: 0.9798 - 2s/epoch
## Epoch 20/30
## 375/375 - 2s - loss: 0.0582 - accuracy: 0.9827 - val_loss: 0.1067 - val_accuracy: 0.9793 - 2s/epoch
## Epoch 21/30
## 375/375 - 2s - loss: 0.0590 - accuracy: 0.9831 - val_loss: 0.1193 - val_accuracy: 0.9781 - 2s/epoch
## Epoch 22/30
## 375/375 - 2s - loss: 0.0566 - accuracy: 0.9835 - val_loss: 0.1095 - val_accuracy: 0.9800 - 2s/epoch
## Epoch 23/30
## 375/375 - 2s - loss: 0.0561 - accuracy: 0.9841 - val_loss: 0.1085 - val_accuracy: 0.9812 - 2s/epoch
## Epoch 24/30
## 375/375 - 2s - loss: 0.0534 - accuracy: 0.9851 - val_loss: 0.1194 - val_accuracy: 0.9794 - 2s/epoch
## Epoch 25/30
## 375/375 - 2s - loss: 0.0561 - accuracy: 0.9846 - val_loss: 0.1098 - val_accuracy: 0.9803 - 2s/epoch
## Epoch 26/30
```

```
## 375/375 - 2s - loss: 0.0517 - accuracy: 0.9860 - val_loss: 0.1206 - val_accuracy: 0.9799 - 2s/epoch
## Epoch 27/30
## 375/375 - 2s - loss: 0.0549 - accuracy: 0.9846 - val_loss: 0.1192 - val_accuracy: 0.9795 - 2s/epoch
## Epoch 28/30
## 375/375 - 2s - loss: 0.0540 - accuracy: 0.9858 - val_loss: 0.1131 - val_accuracy: 0.9808 - 2s/epoch
## Epoch 29/30
## 375/375 - 2s - loss: 0.0486 - accuracy: 0.9858 - val_loss: 0.1221 - val_accuracy: 0.9812 - 2s/epoch
## Epoch 30/30
## 375/375 - 2s - loss: 0.0497 - accuracy: 0.9862 - val_loss: 0.1185 - val_accuracy: 0.9803 - 2s/epoch
```

```
plot(history)
```



```
model %>% evaluate(x_test, y_test)
```


```
## 313/313 - 0s - loss: 0.1102 - accuracy: 0.9816 - 300ms/epoch - 958us/step
##      loss accuracy
## 0.1101681 0.9816000
```

```
model %>% predict(x_test) %>% k_argmax()
```

```
## 313/313 - 0s - 378ms/epoch - 1ms/step
```

```
## tf.Tensor([7 2 1 ... 4 5 6], shape=(10000), dtype=int64)
```

The keras package also plots the progress of the model by showing the loss function evolution by epoch, as well as accuracy, for the training and validation samples.



Forest recommended using PyTorch. There is an R package `torch` that allows you to use PyTorch-like functionality natively from R. No Python installation is required: `torch` is built directly on top of libtorch, a C++ library that provides the tensor-computation and automatic-differentiation capabilities essential to building neural networks. See <https://torch.mlverse.org/> for more information.

Using `nnet` with some

Now we're going to do a Neural Net on some simulated data. We're going to use the `nnet` package for the neural network:

```
library(clusterGeneration)
```

```
## Loading required package: MASS
```

```
library(nnet)
```

Now we'll simulate the data. Part of the difficulty in simulating multivariate normal data is getting a random positive definite matrix. The `clusterGeneration` package helps with this.

```
#define number of variables and observations
set.seed(2)
num.vars<-8
num.obs<-10000

#define correlation matrix for explanatory variables
#define actual parameter values
cov.mat<-genPositiveDefMat(num.vars,covMethod=c("unifcorrmat"))$Sigma
rand.vars<-mvrnorm(num.obs,rep(0,num.vars),Sigma=cov.mat)
parms1<-runif(num.vars,-10,10)
y1<-rand.vars %*% matrix(parms1) + rnorm(num.obs,sd=20)
parms2<-runif(num.vars,-10,10)
y2<-rand.vars %*% matrix(parms2) + rnorm(num.obs,sd=20)

#prep data and create neural network
rand.vars<-data.frame(rand.vars)
resp<-apply(cbind(y1,y2),2, function(y) (y-min(y))/(max(y)-min(y)))
resp<-data.frame(resp)
names(resp)<-c('Y1','Y2')
```

This model has 8 variables with 2 outcomes and a total of 10,000 observations. We'll fit this model with 8

hidden levels with the **nnet** function.

Some info about the **nnet** function:

- it is used for single-hidden layer networks only.
- The call is:

```
nnet(x, y, weights, size, Wts, mask, linout = FALSE, entropy = FALSE, softmax = FALSE,
censored = FALSE, skip = FALSE, rang = 0.7, decay = 0, maxit = 100, Hess = FALSE, trace =
TRUE, MaxNWts = 1000, abstol = 1.0e-4, reltol = 1.0e-8, ...)
```
- Some of the options (from the help file):
 - formula-A formula of the form `class ~ x1 + x2 + ...`
 - x-matrix or data frame of x values for examples.
 - y-matrix or data frame of target values for examples.
 - weights-(case) weights for each example – if missing defaults to 1.
 - size-number of units in the hidden layer. Can be zero if there are skip-layer units.
 - Wts-initial parameter vector. If missing chosen at random.
 - linout-switch for linear output units. Default logistic output units.
 - entropy-switch for entropy (= maximum conditional likelihood) fitting. Default by least-squares.
 - softmax-switch for softmax (log-linear model) and maximum conditional likelihood fitting. `linout`, `entropy`, `softmax` and `censored` are mutually exclusive. `censored` A variant on softmax, in which non-zero targets mean possible classes. Thus for softmax a row of (0, 1, 1) means one example each of classes 2 and 3, but for censored it means one example whose class is only known to be 2 or 3.
 - rang-Initial random weights on `[-rang, rang]`. Value about 0.5 unless the inputs are large, in which case it should be chosen so that `rang * max(|x|)` is about 1.
 - decay-parameter for weight decay. Default 0.
 - maxit-maximum number of iterations. Default 100.
 - abstol-Stop if the fit criterion falls below `abstol`, indicating an essentially perfect fit.
 - reltol-Stop if the optimizer is unable to reduce the fit criterion by a factor of at least `1 - reltol`.


```
mod1<-nnet(rand.vars,resp,size=8,linout=T)
```

```
## # weights: 90
## initial value 30066.949169
## iter 10 value 131.361932
## iter 20 value 51.449404
## iter 30 value 44.110444
## iter 40 value 39.959512
## iter 50 value 37.678074
## iter 60 value 35.715896
## iter 70 value 34.182676
## iter 80 value 33.162657
## iter 90 value 32.665243
## iter 100 value 32.413667
## final value 32.413667
## stopped after 100 iterations
```

To plot the network we're going to import a function from Github

```
library(devtools)
library(reshape)
source_url('https://gist.github.com/fawda123/7471137/raw/cd6e6a0b0bdb4e065c597e52165e5ac887f5fe95/nnet_
et_plot_update.r')

plot.nnet(mod1)
```



Full url is 'https://gist.github.com/fawda123/7471137/raw/cd6e6a0b0bdb4e065c597e52165e5ac887f5fe95/nnet_et_plot_update.r'

The shows all of the connections that fit. The black (gray) lines indicate a positive (negative) association. The thickness of the line indicates the strength of the association.


We can also visualize the effect that each variable has on the outcome using a method discribed in Gevrey et al. (2003) [1] (see below for reference). The hard thing about visualizing the effect one variable, say X_1 , has on the outcome, say Y , is that all of the other variables, say X_2, X_3, \dots, X_r interact with this relationship. As a result, we visualize the effect between X_1 and Y for fixed values of X_2, X_3, \dots, X_r . One way to do this is to:

- Set X_2, X_3, \dots, X_r to a fixed quantile, for example Q_1 , then feed all values of X_1 into the network and calculate the expected value of Y for each value.
- Plot the expectation of Y by X_1 .
- Set X_2, X_3, \dots, X_r to a different fixed quantile, repeat as necessary.


This method will give you some idea of what is impacting the outcome, but it will not give any information on the interactions between the variables.

```
library(ggplot2)
source('https://gist.githubusercontent.com/fawda123/6860630/raw/b8bf4a6c88d6b392b1bfa6ef24759ae98f31877
lek.fun(mod1)
```

```
## Warning: Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.
## i Please use `linewidth` instead.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```



```
lek.fun(mod1,var.sens=c('X2','X5'),split.vals=seq(0,1,by=0.05))
```

Full url is 'https://gist.githubusercontent.com/fawda123/6860630/raw/b8bf4a6c88d6b392b1bfa6ef24759ae98f31877c/lek_fun.r'


Next we're going to look at some measures of variable importance. The goal here is to measure the relative importance of explanatory variables for their relation to one or more response variables. The method we'll use here is described in Garson (1991) [2]. Here is the basic idea:

- Identify all weighted connections between a specific input node and a specific response variable.
- Do this for all other explanatory variables until we have a list of all weights for each input variable.
- The connections are "tallied" for each input node and scaled relative to all other inputs.
- A single value is obtained for each explanatory variable that describes the relationship with response variable in the model (see the appendix in Goh 1995 [3] for a more detailed description).


To do this we will import 'gar.fun' from Github (this is not a package just a function).

```
source_url('https://gist.githubusercontent.com/fawda123/6206737/raw/d6f365c283a8cae23fb20892dc223bc5764')
#create a pretty color vector for the bar plot
cols<-colorRampPalette(c('lightgreen', 'lightblue'))(num.vars)

#use the function on the model created above
par(mar=c(3,4,1,1),family='serif')
gar.fun('Y1',mod1)
```



```
gar.fun('Y2',mod1)
```



- [1] Gevrey M, Dimopoulos I, Lek S. 2003. Review and comparison of methods to study the contribution of variables in artificial neural network models. *Ecological Modelling*. 160:249-264.
- [2] Garson, G.D. 1991. Interpreting neural network connection weights. *Artificial Intelligence Expert*. 6(4):46–51.
- [3] Goh, A.T.C. 1995. Back-propagation neural networks for modeling complex systems. *Artificial Intelligence in Engineering*. 9(3):143–151.